

Knative Serverless and Kubernetes

Consistent automated software deployments with Knative

Andrew Ripka
@rippmn

Context Setting

Photo by [Michael Kogan](#), Licensed under CC BY-SA 2.0

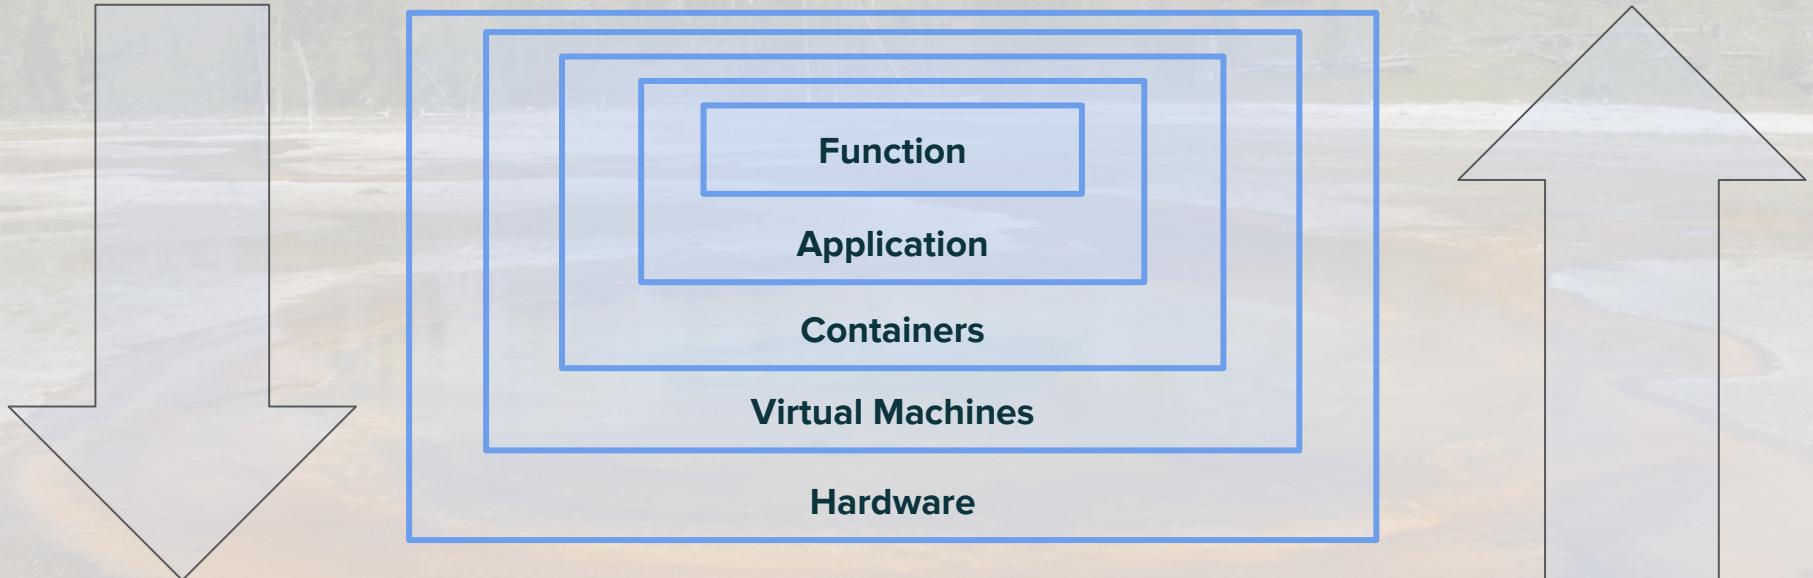


“Serverless”
Running a workload without worrying
about....

Platform Developer Abstractions

More Standards

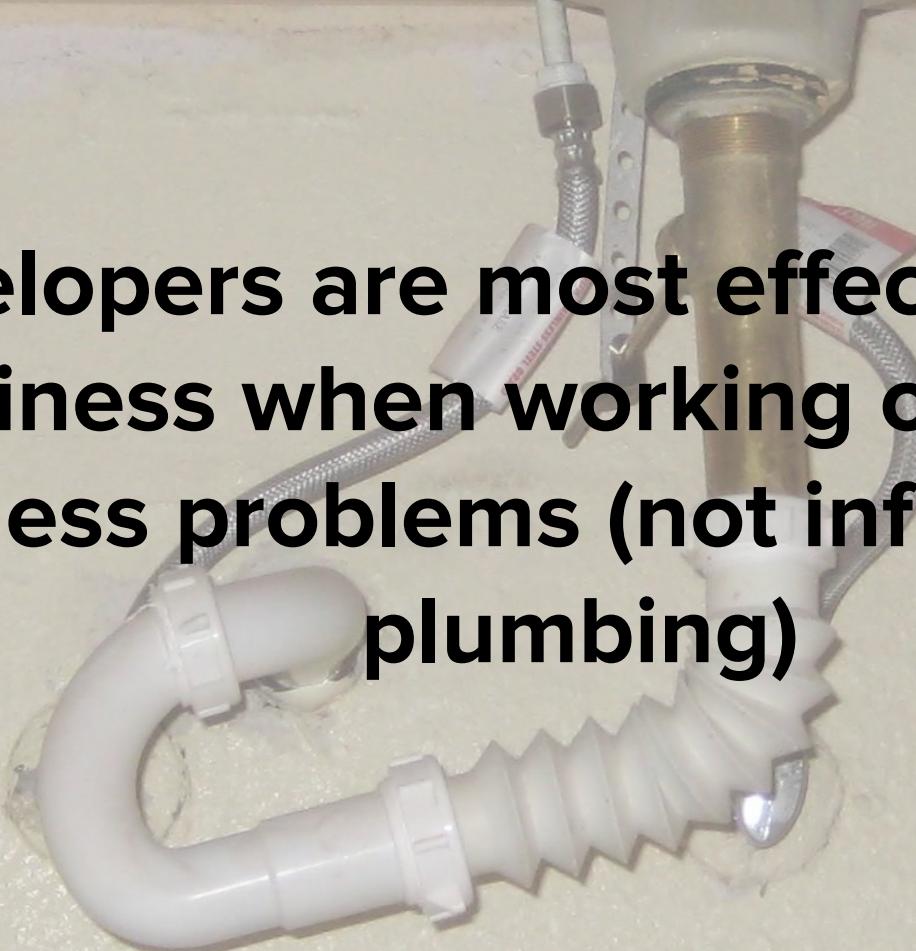
Higher Efficiency



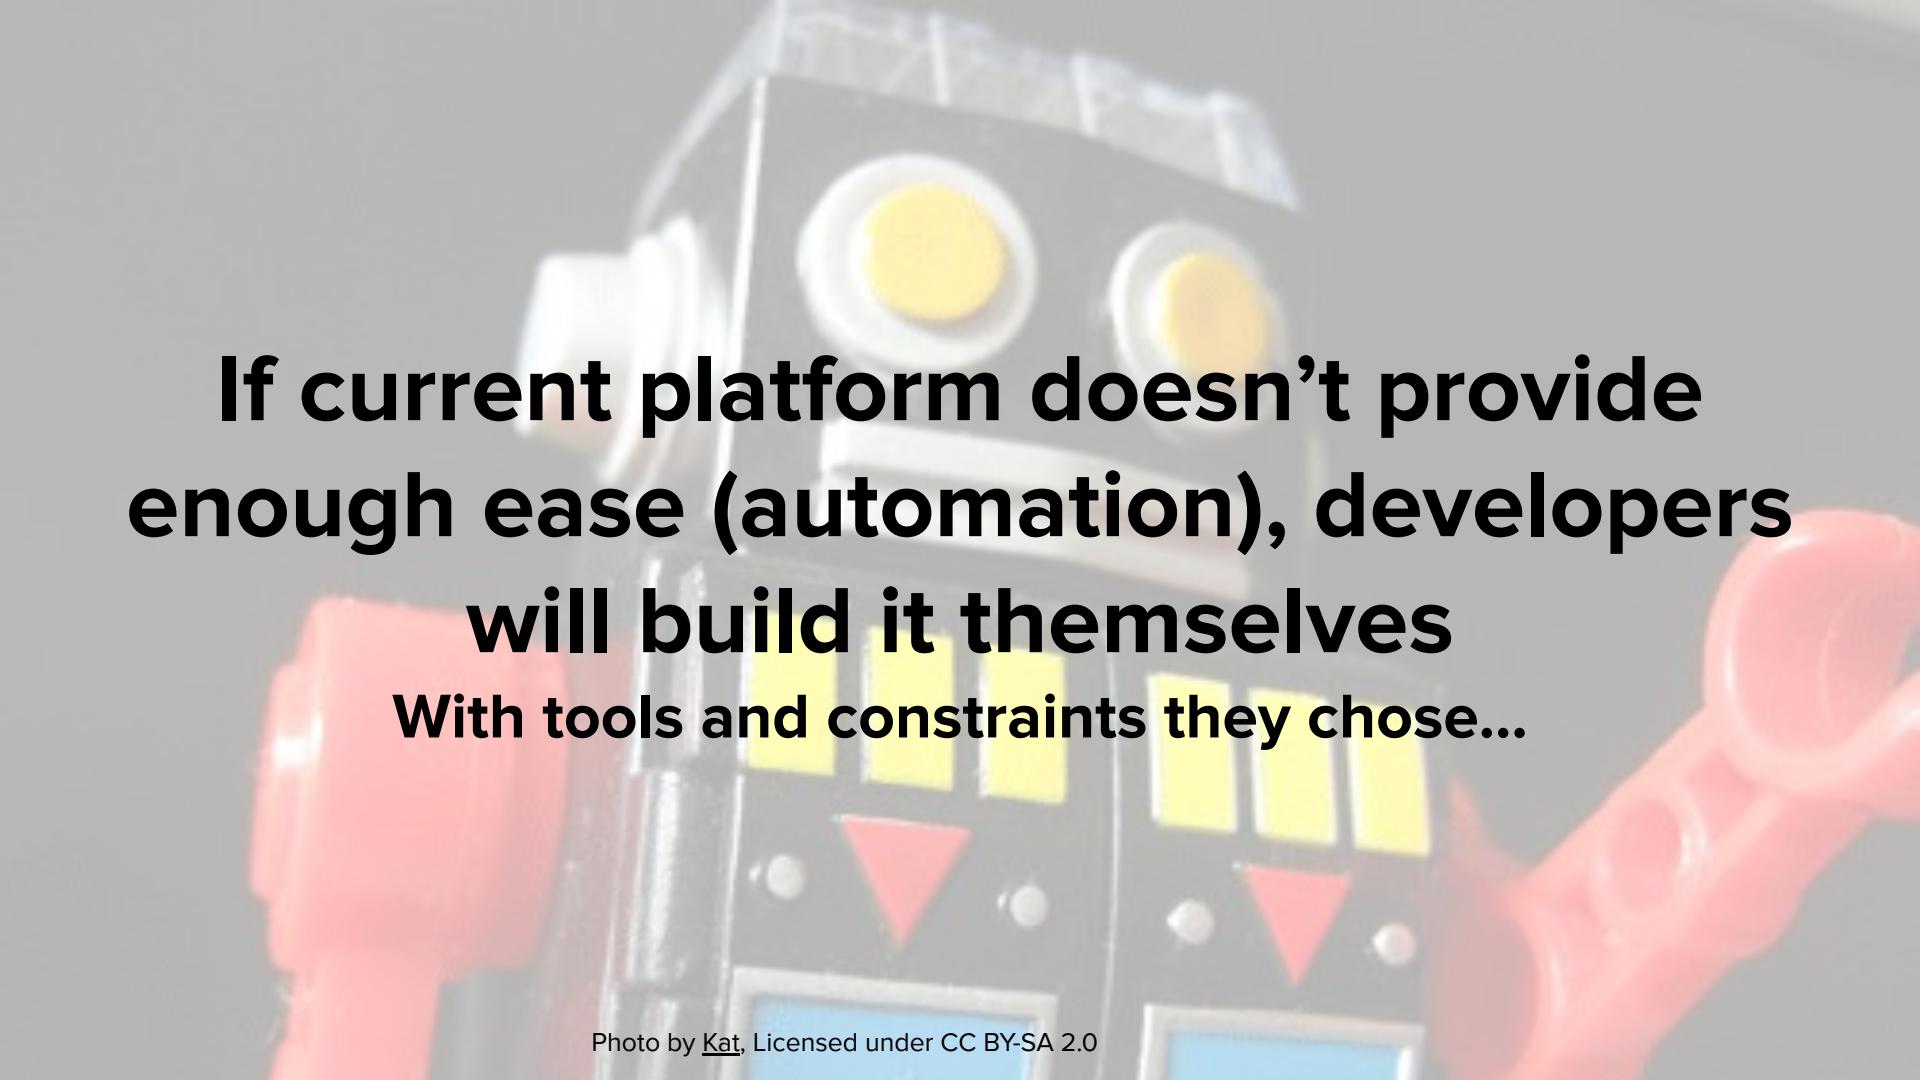
More Flexibility

What is Serverless?

- Frees devs from details that aren't part of the core business logic of their code. (e.g. Infrastructure)
- Small single purpose functions that are invoked in response to events
- Scale to 0 when no events occurring



Developers are most effective to the business when working on solving business problems (not infrastructure plumbing)



**If current platform doesn't provide
enough ease (automation), developers
will build it themselves
With tools and constraints they chose...**



**Ad hoc automation is awesome, until
someone who wasn't involved gets
paged.
(Like the operations team)**

**A platform's overall goal should be to
make the right thing to do the easiest
thing to do
(thus providing automation with desired constraints)**

A photograph of a stage setup. On the left, there are two Marshall brand amplifiers. In the center, a drum set is positioned. To the right, a guitar is leaning against a microphone stand. The stage is illuminated by several spotlights hanging from the ceiling, casting a warm glow on the equipment.

So you say you want
a serverless platform

and you want to do it yourself...



Kelsey Hightower 
@kelseyhightower

I'm convinced the majority of people managing infrastructure just want a PaaS. The only requirement: it has to be built by them.

729 6:08 PM - Apr 11, 2017

388 people are talking about this >



Why not Kubernetes?

It's building blocks



Kelsey Hightower

@kelseyhightower



Kubernetes is a platform for building platforms. It's a better place to start; not the endgame.

635 4:04 PM - Nov 27, 2017



226 people are talking about this





An infrastructure abstraction



Dan Woods
@danveloper

3. You are building an application platform. A lot of people treat Kubernetes _as_ the platform, but it provides the foundational models for infrastructure, not applications. Higher-order modeling is required to translate the infra domain to the app domain.

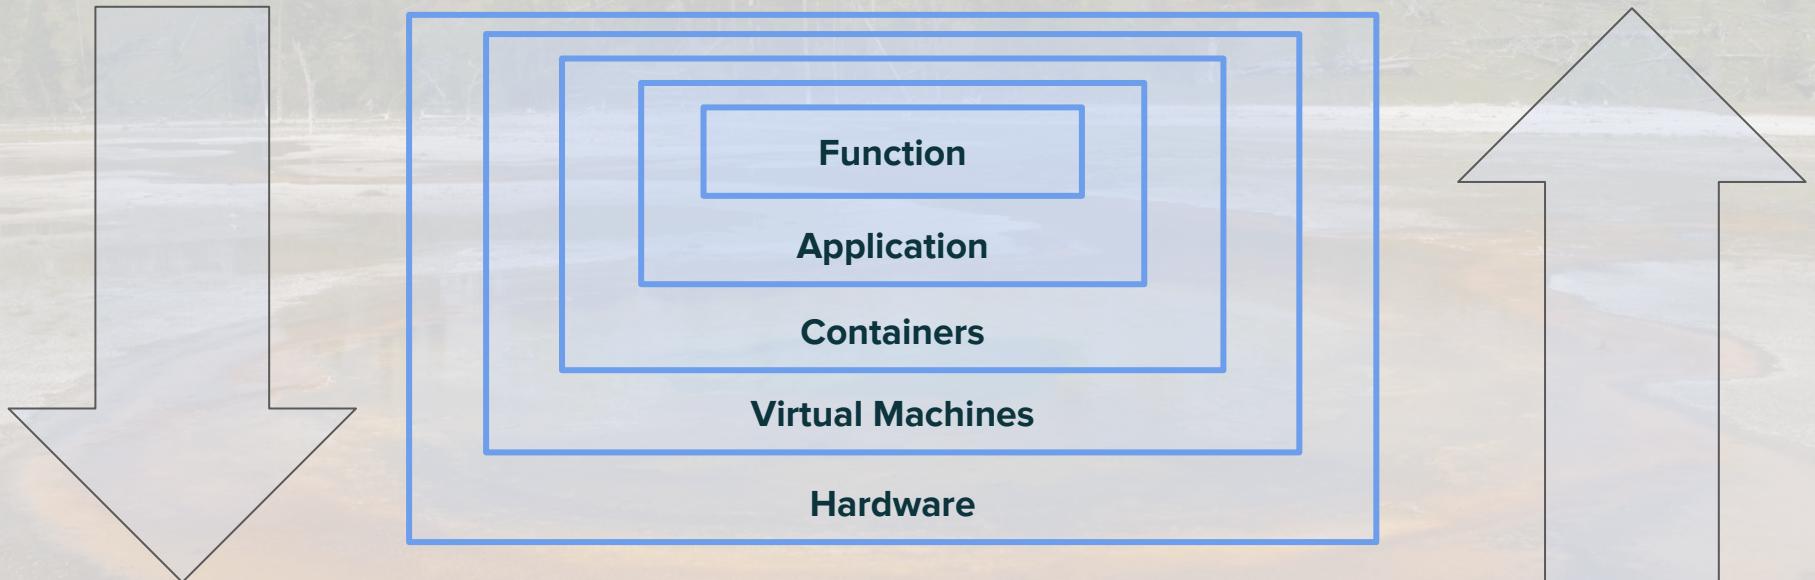
♡ 55 6:58 AM - Apr 19, 2019



Platform Developer Abstractions

More Standards

Higher Efficiency



More Flexibility

Serverless Platform Coarse Grained Requirements

- Package to deployable
- Deploy and Run
- Map traffic and events to running handlers

Serverless Platform

- Contains container creation and deployment process automation
- Allows scale to zero and back up when necessary
- Handling of revisions with rollback with traffic splitting
- An eventing system with configurable sources/flows/subscribers
- Something to collect metrics and export telemetry from the app.



Knative

What is Knative is not

- It's not a fully featured Function as a Service platform
- It's intended to be used by others to build FaaS or PaaS products

Kelsey Hightower 
@kelseyhightower

Following ▾

Replies to @cpuguy83

Knative is not that top-level interface. That will come from the ecosystem. The containers and specs that describes how workloads run and flow between those systems is the value of Knative.

7:29 AM - 8 Mar 2019

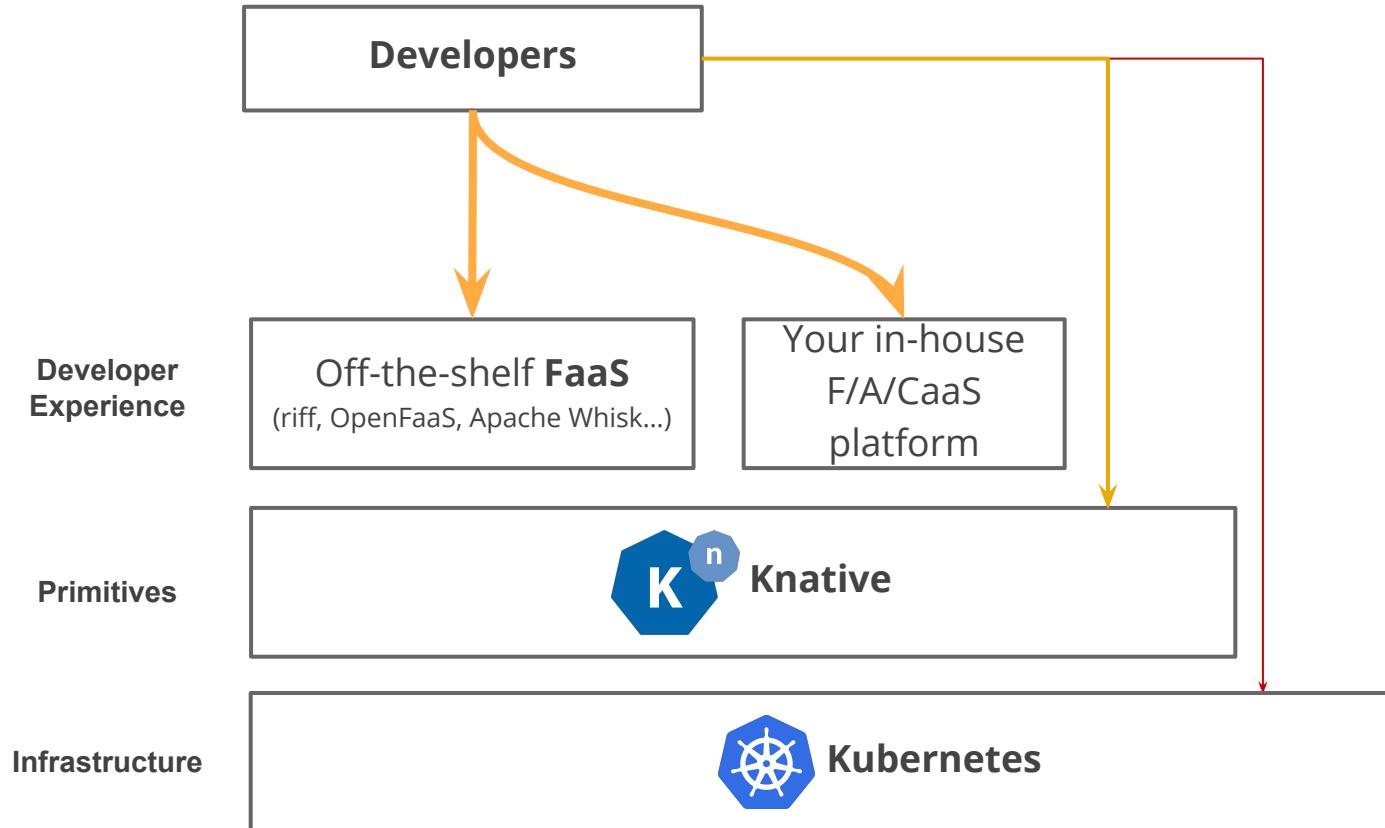
What is Knative

- An abstraction on top of Kubernetes
 - Deploys and runs containers
- Set of building blocks to construct your own FaaS/AaaS/CaaS
 - abstracts common tasks through custom Kubernetes API objects
- Set of middleware components essential to build modern, source-centric, and container-based applications that can run anywhere.

Knative Codified Patterns

- Deploying a container
- Orchestrating source-to-URL workflows on Kubernetes
- Routing and managing traffic with blue/green deployment
- Scaling automatically and sizing workloads based on demand
- Binding running services to eventing ecosystems

Expected usage



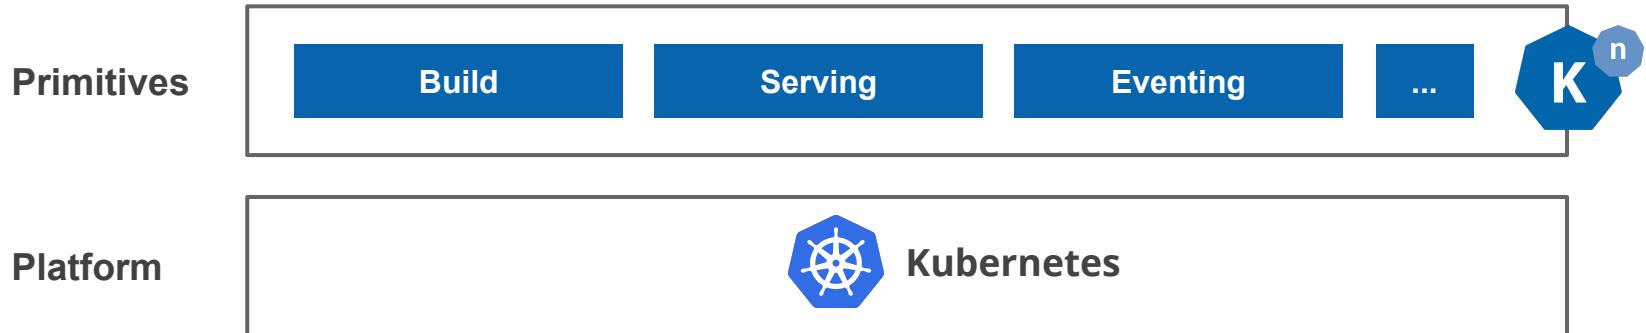
Knative architecture

Platform

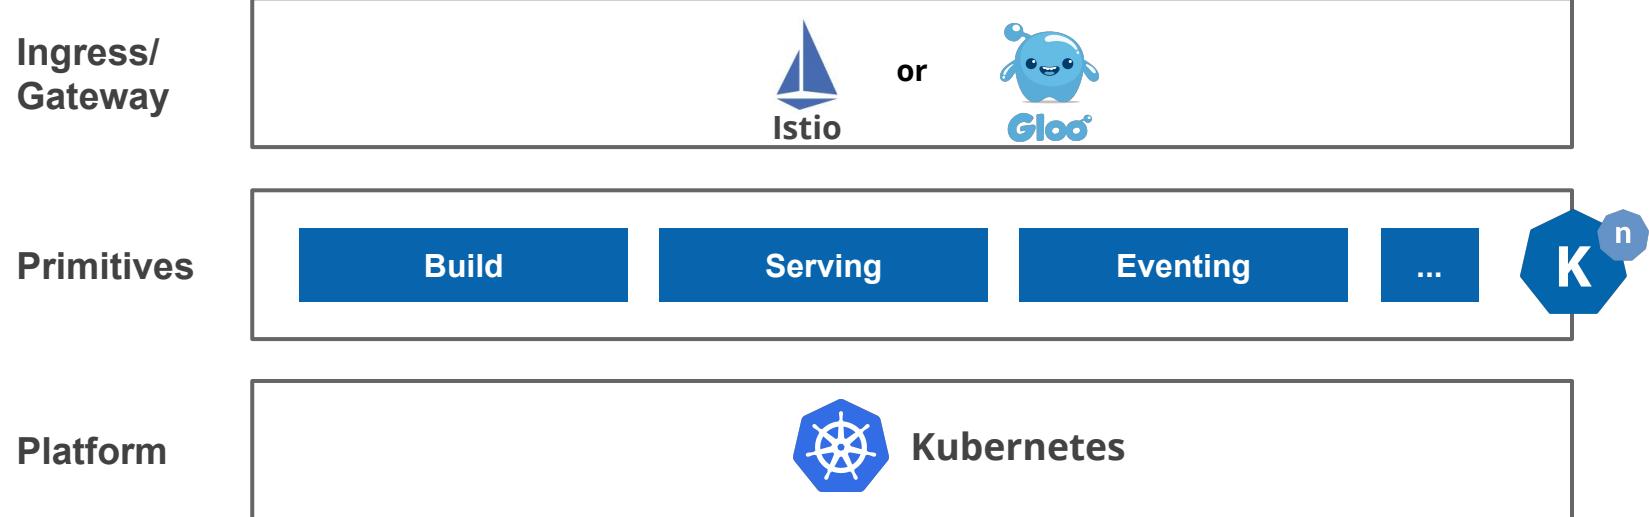


Kubernetes

Knative architecture



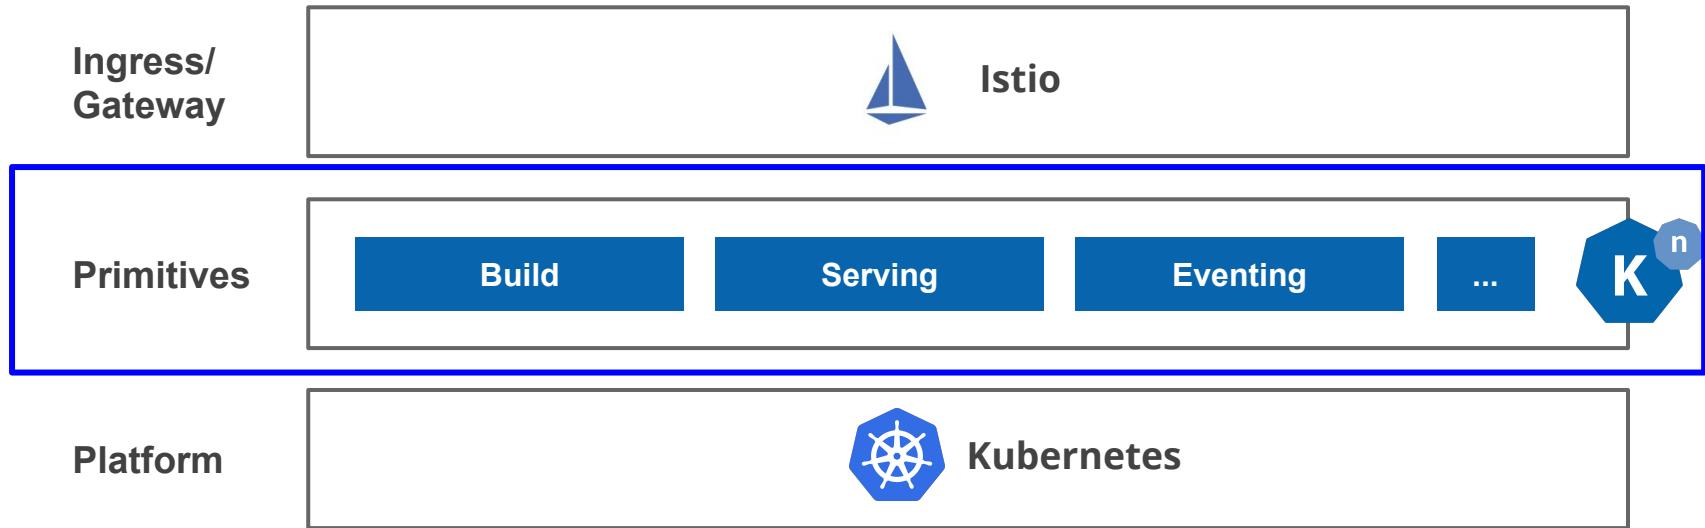
Knative architecture



Audiences



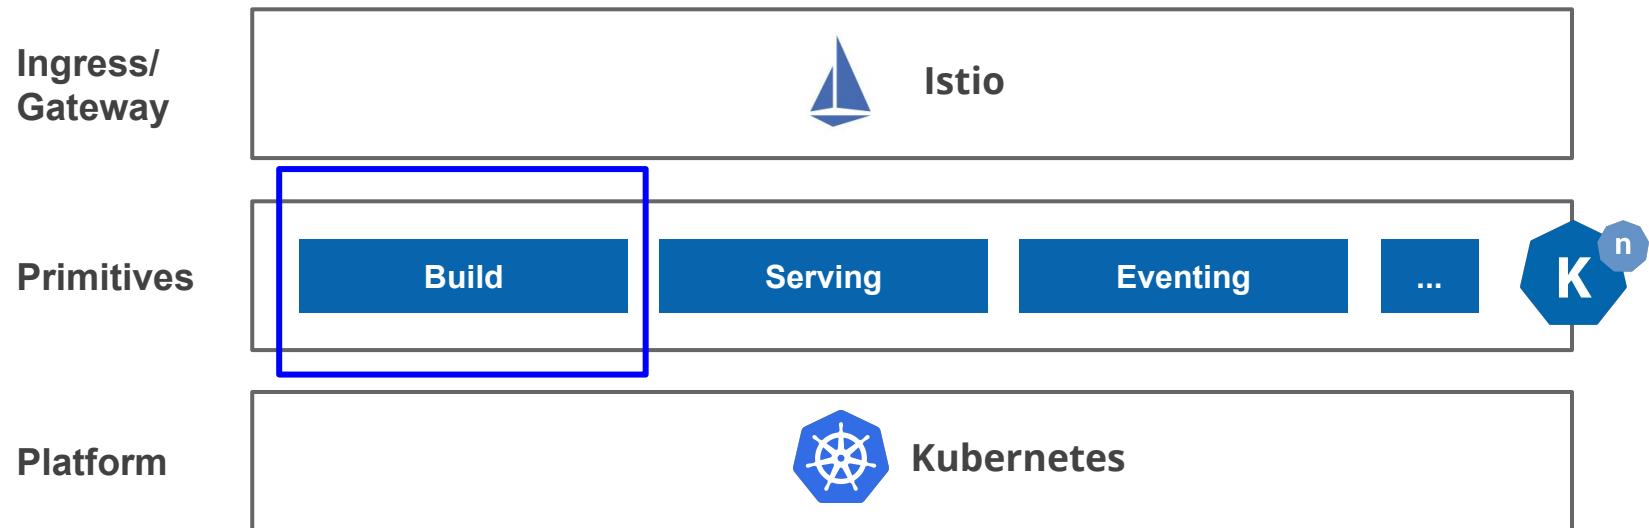
Knative Components



Knative Components

- **Build** - Source-to-container build orchestration
- **Serving** - Request-driven serverless compute using containers with revisions and scale to zero
- **Eventing** - Declarative management and delivery of events

Knative Components



Build

- Defines process that runs to completion and can provide status.
- Build container image from source that can be deployed to Knative Serving
- Uses Kubernetes Service Account to connect to secured resources

Build Components

- **Build** is made up of Steps containing builders
- **Builder** - container image that accomplishes a task (discrete or composite)
 - One of the steps typically push to container registry
- **Source** - data necessary for build in mounted volume
 - GCS
 - Git
 - Custom (container image)

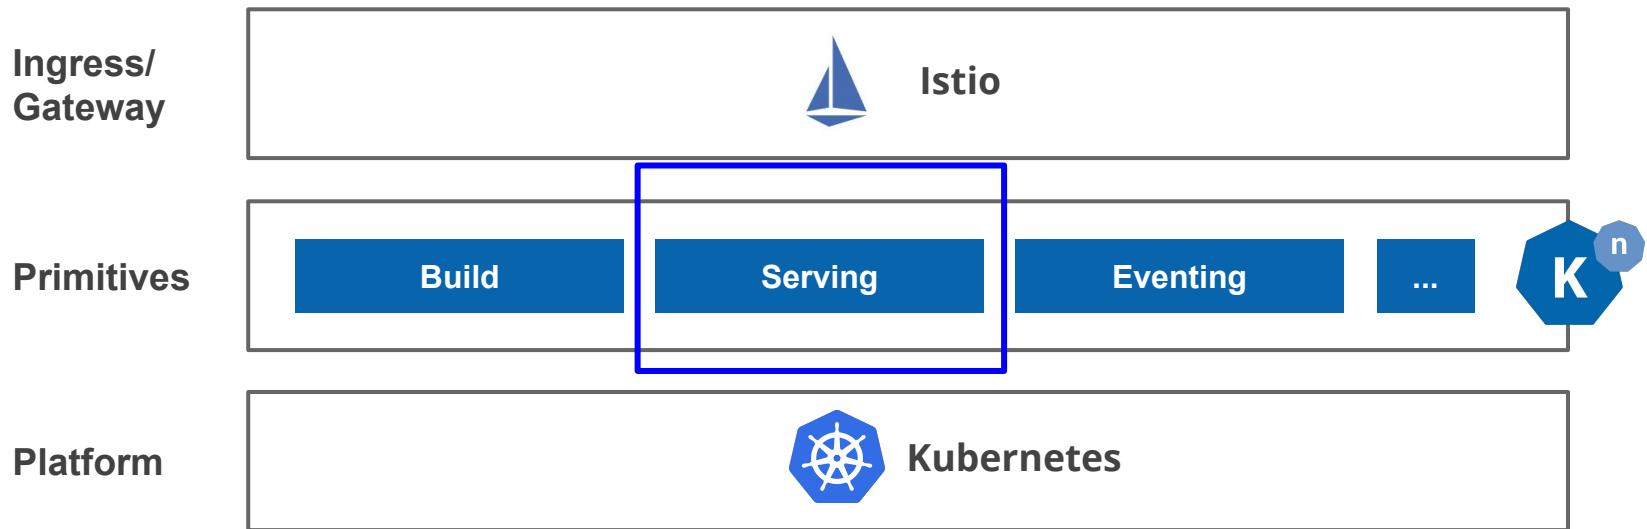
Build Components (Cont.)

- **Service Account** - privileges to run build with
- **Volumes** - non default volumes to use during build steps
- **Timeout** - max time for resource allocation and build steps

Build Templates

- Encapsulate build process for reuse and consistency
 - limited parameterization
- Similar to “compile” phase of Heroku and Cloud Foundry buildpacks
- Limited to namespace but can be shared across cluster as Cluster Build Template

Knative Components

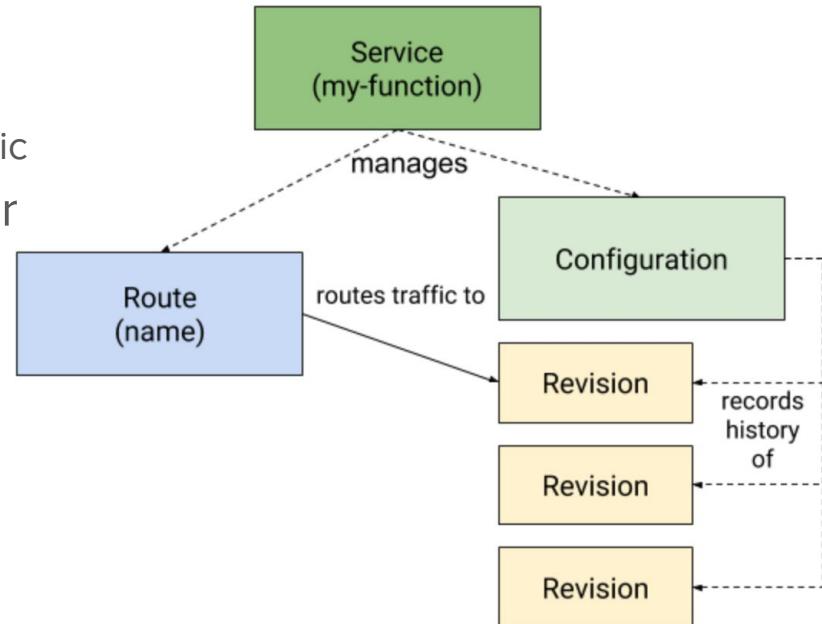


Serving

- Rapid deployment of serverless containers
- Automatic scaling up and down to zero
- Routing and network programming for ingress gateway components (Istio or Gloo)
- Point-in-time snapshots of deployed code and configurations (enabling rollback)
- Provide observability through logs and metrics

Serving Components

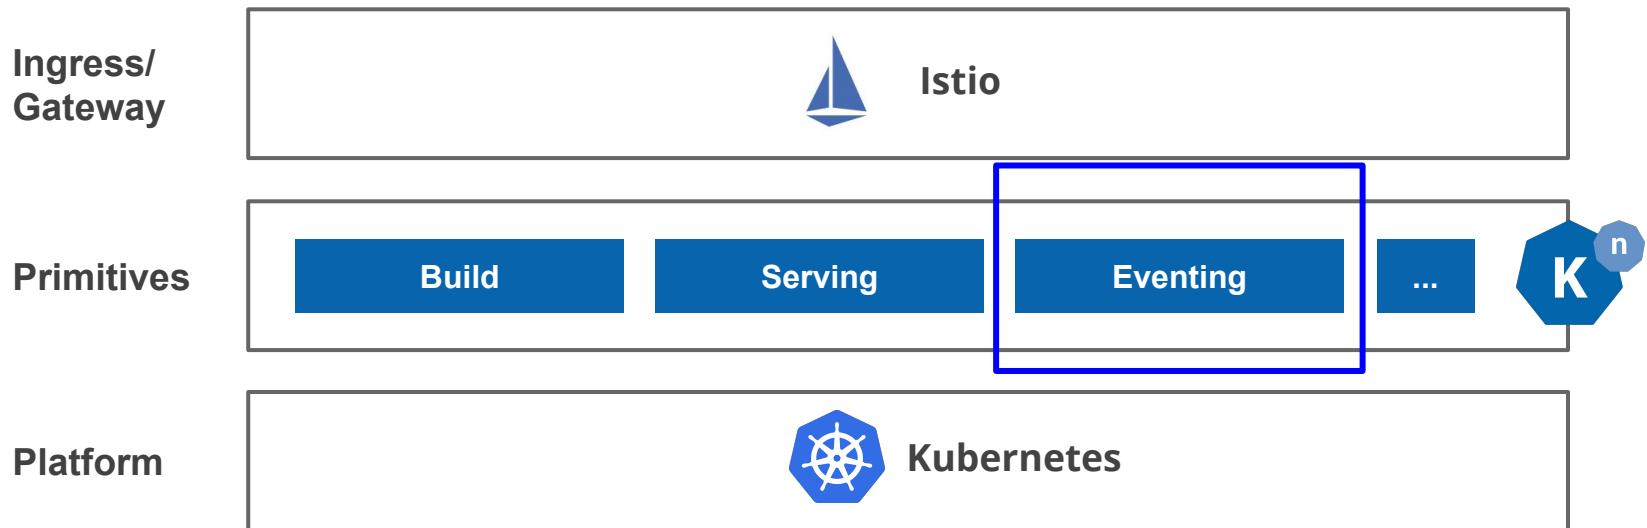
- **Route** map Http routable endpoint to revision(s)
 - Allows for percentage based splitting of traffic
- **Revision** - immutable code (as container or build) and config
 - Generated by configuration changes
 - Retried no route mapped to (resources deleted) but saved for rollback
- **Configuration** - Latest revision of code and config
- **Service** - Set of routes and configurations
 - Single abstraction of software component



Logging And Metrics

- Knative collects logs using Fluentd
- Pluggable logging observability implementations
 - Elasticsearch and Kibana
 - Stackdriver
 - Custom logging plugin
- Provides Grafana metrics observability

Knative Components

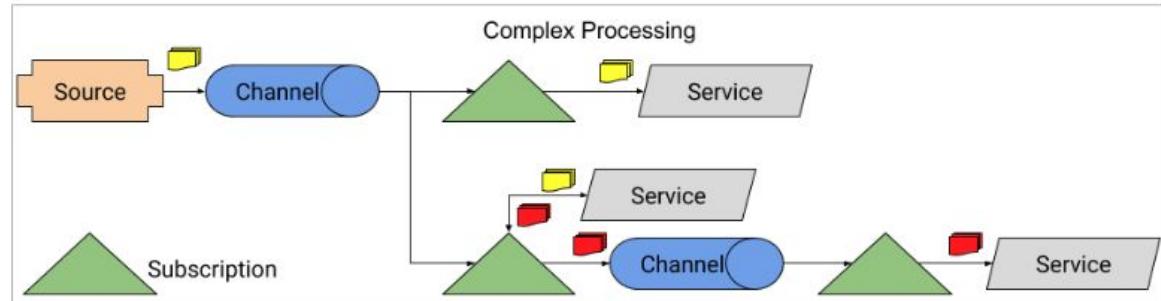


Eventing

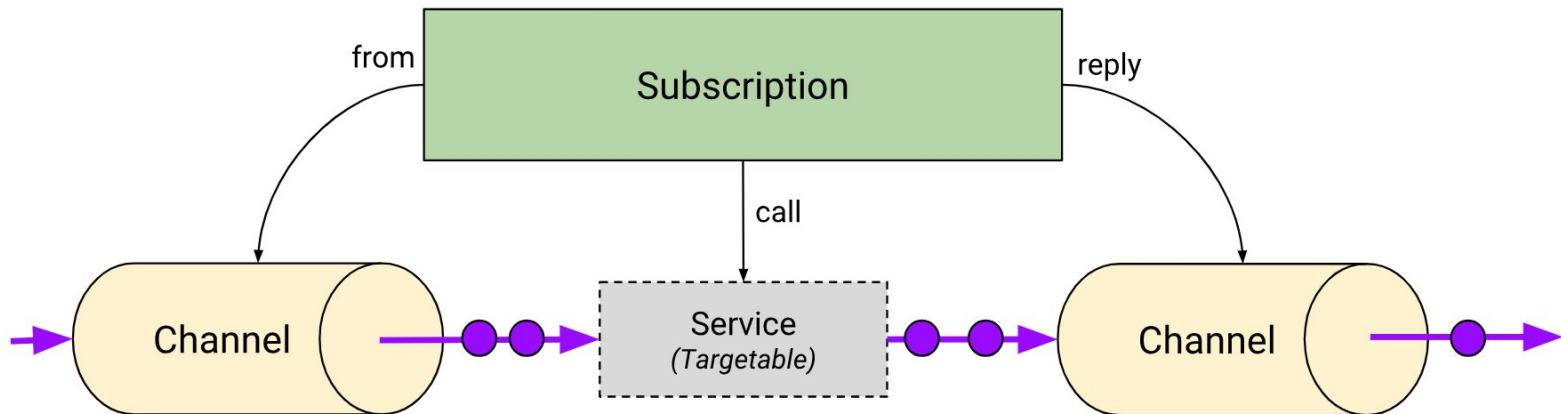
- Composable primitives enabling late-binding loosely coupled event producers and event consumers
- Consumers can be other Services allowed to:
 - Create new applications without modifying the event producer or event consumer.
 - Select and target specific subsets of the events from their producers.
- Consistent with CloudEvents specification developed by CNCF Serverless WG

Eventing Components (Legacy)

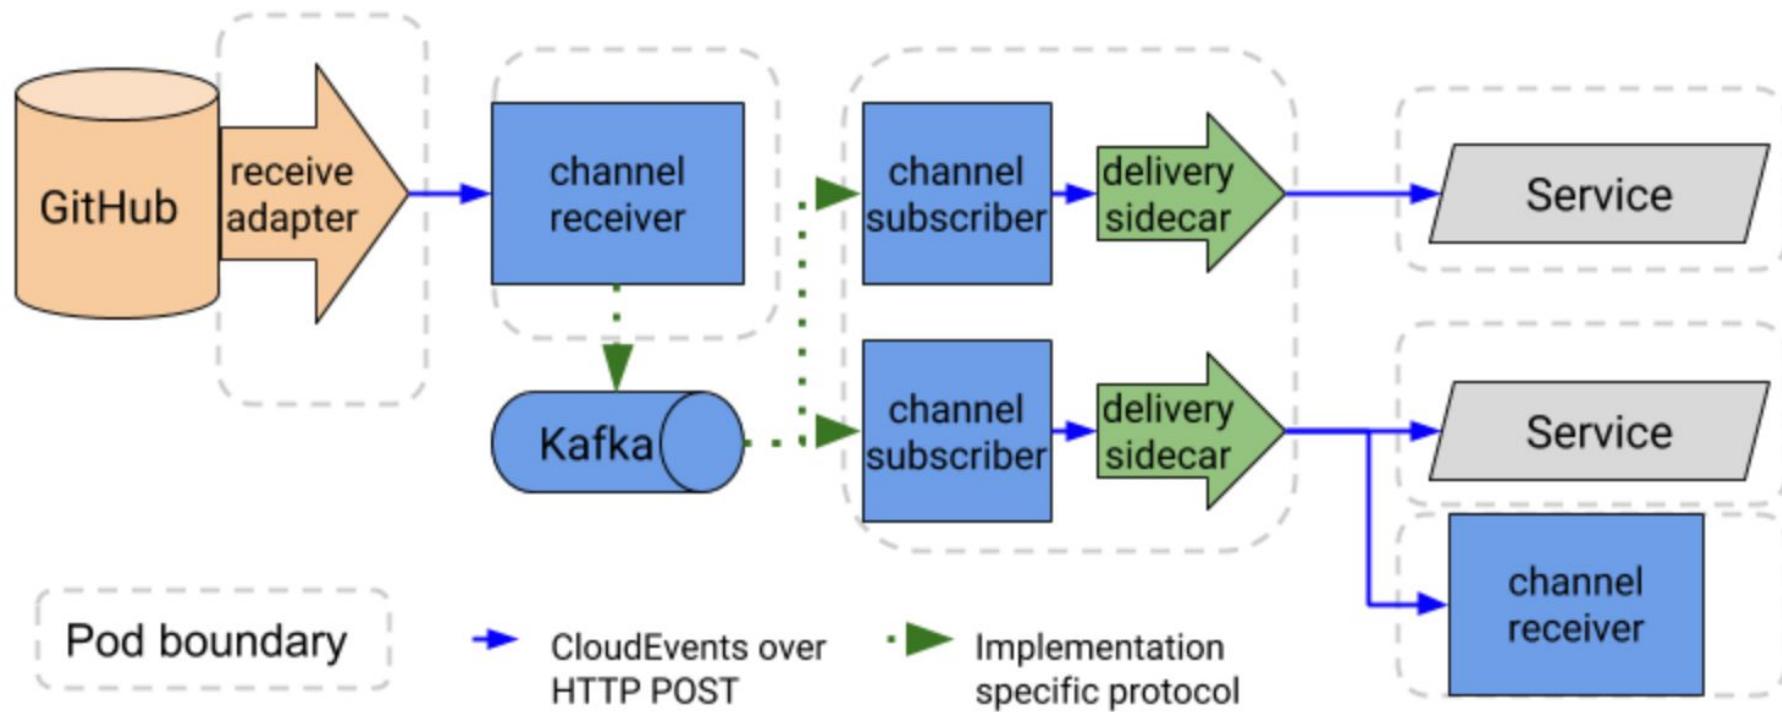
- **Source** - Register interest in events from other systems and publish to a sink
 - K8s events, Kafka, GitHub, GitLab, BitBucket, custom sources
- **Service** - Consumer of events
 - Addressable acknowledge delivery
 - Callable - return 0 or 1 events
- **Channel** - event forwarding persistence
 - Pluggable ClusterChannelProvisioner implementation for message providers
- **Subscription** - means to declare interest in events from a channel



Channels & Subscriptions

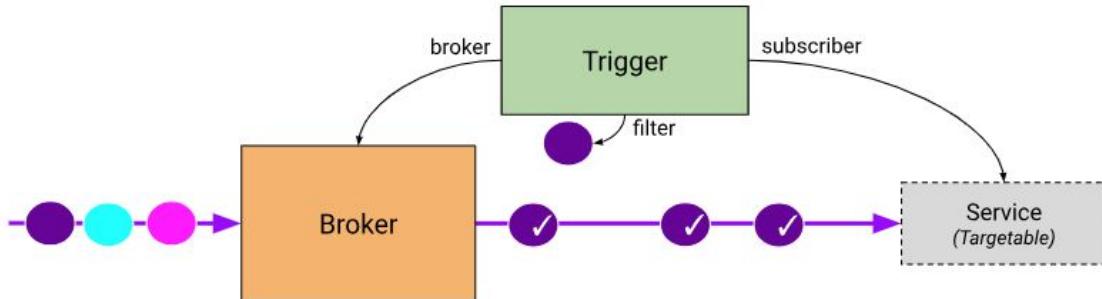


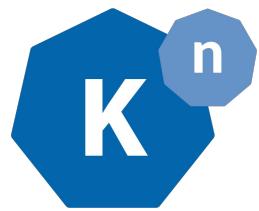
Data Plane



Eventing Components (New as of Knative 0.5)

- **Broker** - mesh of events sent to interested subscribers
 - Use a channel template to create internal channels
- **Trigger** - declares subscription to events from a broker
 - Can be used to filter events based upon type





Config Samples

Config Sample Links

- Build
 - Build - <https://github.com/GoogleCloudPlatform/knative-build-tutorials/blob/master/docker-build/README.md>
 - Build Template - <https://github.com/knative/build-templates/blob/master/kaniko/kaniko.yaml>
 - Build with template - <https://github.com/knative/build-templates/tree/master/kaniko>
- Serving
 - Deploy and routing examples -
<https://github.com/knative/docs/blob/master/docs/serving/samples/blue-green-deployment.md>
- Eventing
 - Channel - <https://github.com/gswk/knative-eventing-demo/blob/master/channel.yaml>
 - Subscription - <https://github.com/gswk/knative-eventing-demo/blob/master/subscription.yaml>
 - Broker - <https://github.com/knative/eventing/tree/master/docs/broker>
 - Trigger - <https://github.com/knative/docs/blob/master/docs/eventing/samples/gcp-pubsub-source/trigger.yaml>

A large stack of red and tan clay bricks, some stacked vertically and others horizontally, creating a textured, layered appearance.

**Remember these are primitives to build
PaaS/FaaS**

Where to Go to learn more

- <http://knative.dev>
- <https://github.com/knative/>
- <https://starkandwayne.com/blog/introduction-to-knative/>
- <https://medium.com/@pczarkowski/introduction-to-knative-b93a0b9aeeef>
- <https://www.ibm.com/cloud/learn/knative>
- <https://cloud.google.com/knative/>