

Knative

Bringing Serverless to K8s

Andrew Ripka
@rippmn

Context Setting

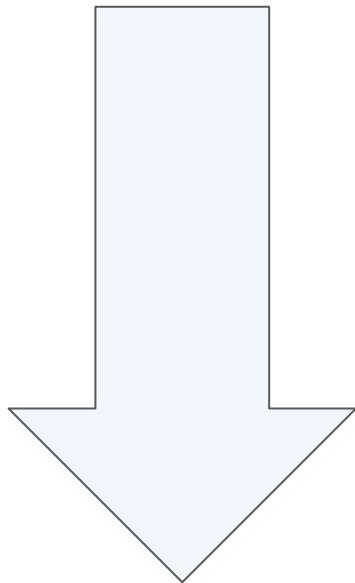
Photo by [Michael Kogan](#), Licensed under CC BY-SA 2.0



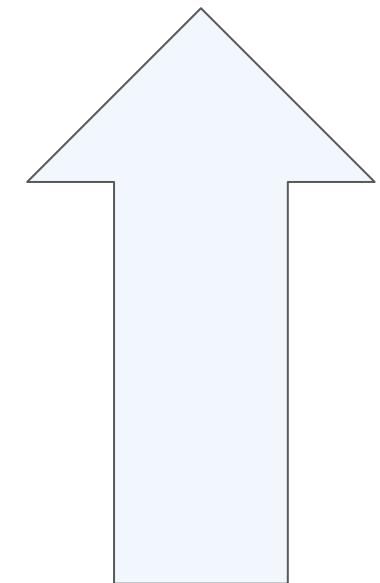
“Serverless”
Running a workload without
worrying about....

Platform Developer Abstractions

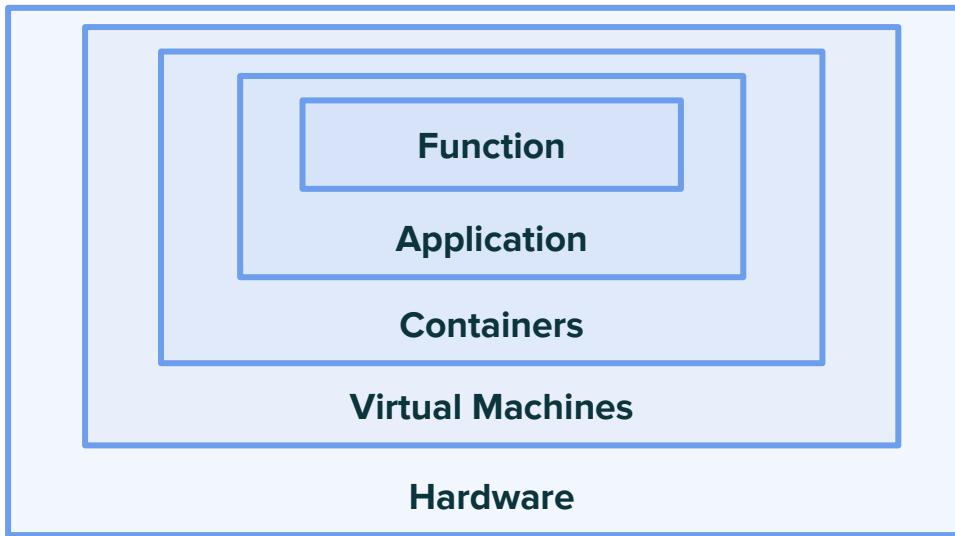
More Standards



Higher Efficiency



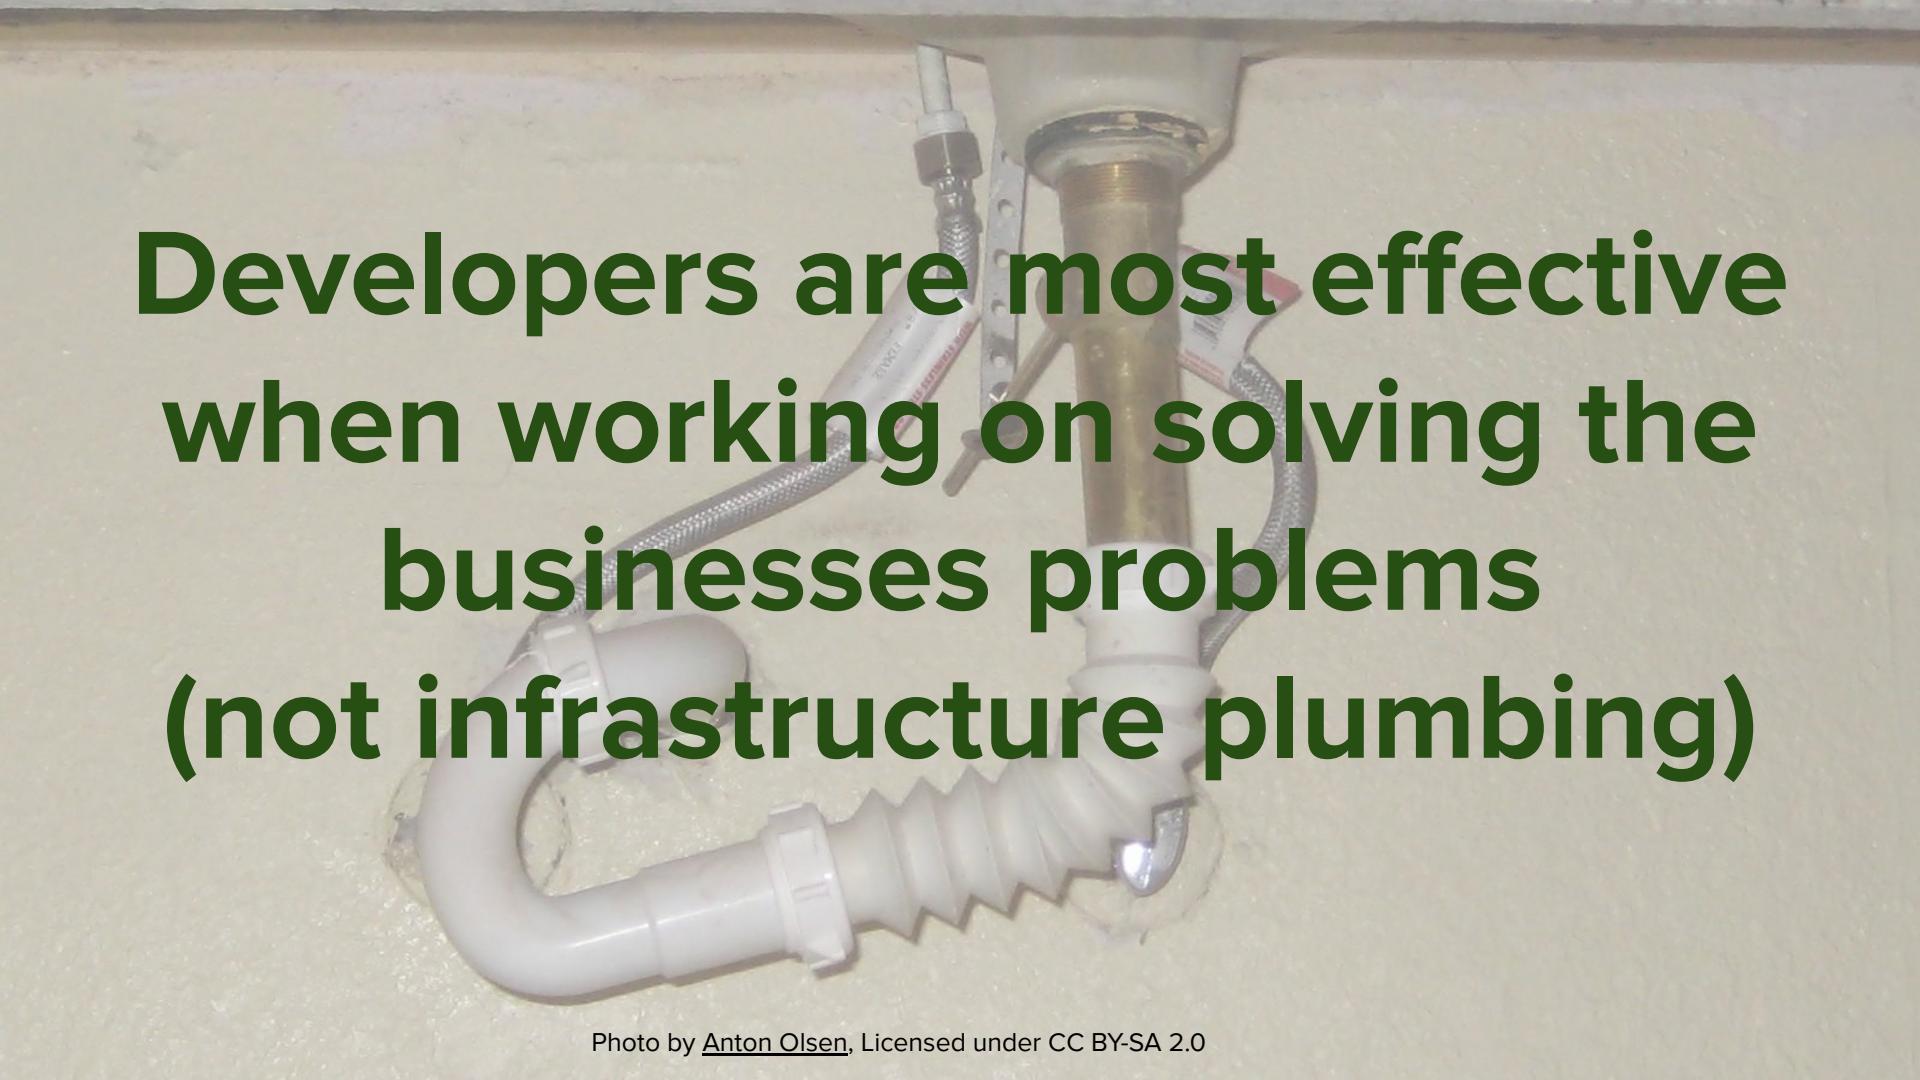
More Flexibility



Lower Efficiency

What is Serverless?

- **Frees devs from details that aren't part of the core business logic of their code. (e.g. Infrastructure)**
- **Small single purpose functions that are invoked in response to events**
- **Scale to 0 when no events occurring**



**Developers are most effective
when working on solving the
businesses problems
(not infrastructure plumbing)**

A grayscale photograph showing a close-up of a person's hands working on a dark-colored electronic circuit board. Several yellow sticky notes are attached to the board, some with handwritten text like 'VDD' and 'GND'. The background is blurred, suggesting a workshop or laboratory environment.

If current
platform isn't easy enough
developers
will build one themselves
With tools and constraints they chose...

**Ad hoc automation is awesome,
until someone who wasn't involved
gets paged.
(Like the operations team)**

A photograph showing a person's silhouette cast onto a light-colored wall or surface. The shadow is elongated and slightly curved, with the words "make the right thing to do" faintly visible where the person's body blocks the light.

A platform's overall goal
make the right thing to do
the easiest thing to do
(automation providing desired
constraints)

A photograph of a stage setup. In the foreground, there are two Marshall brand guitar amplifiers. Behind them is a drum set with cymbals. The stage is illuminated by several spotlights hanging from the ceiling, casting a warm glow. The background is dark, suggesting a concert or performance setting.

So you say you want a serverless platform

and you want to do it yourself...



Kelsey Hightower

@kelseyhightower



I'm convinced the majority of people managing infrastructure just want a PaaS. The only requirement: it has to be built by them.

729 6:08 PM - Apr 11, 2017



388 people are talking about this





Why not Kubernetes?



Kelsey Hightower



@kelseyhightower



Kubernetes is a platform for building platforms. It's a better place to start; not the endgame.

635 4:04 PM - Nov 27, 2017



226 people are talking about this



It's building blocks

A photograph of the interior of a large, curved concrete tunnel. The ceiling and walls are made of light-colored concrete and feature various pipes, cables, and structural supports. The perspective is looking down the length of the tunnel.

And an infrastructure abstraction



Dan Woods
@danveloper

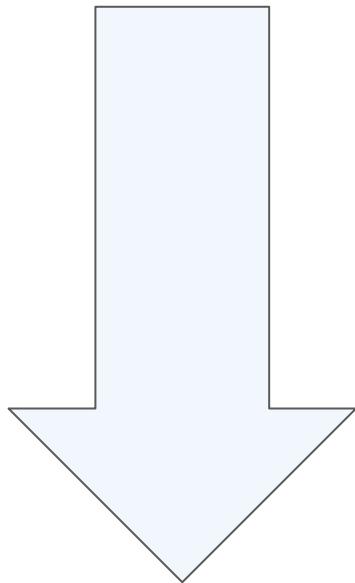
3. You are building an application platform. A lot of people treat Kubernetes _as_ the platform, but it provides the foundational models for infrastructure, not applications. Higher-order modeling is required to translate the infra domain to the app domain.

♡ 55 6:58 AM - Apr 19, 2019

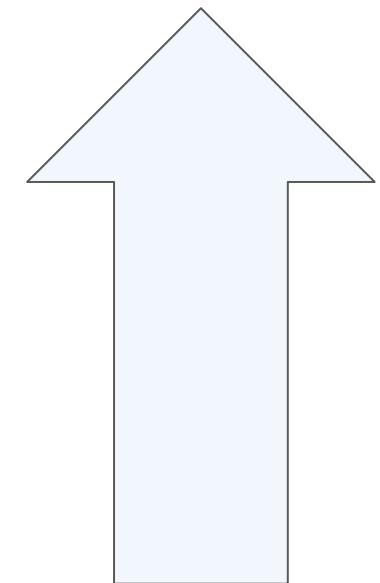


Platform Developer Abstractions

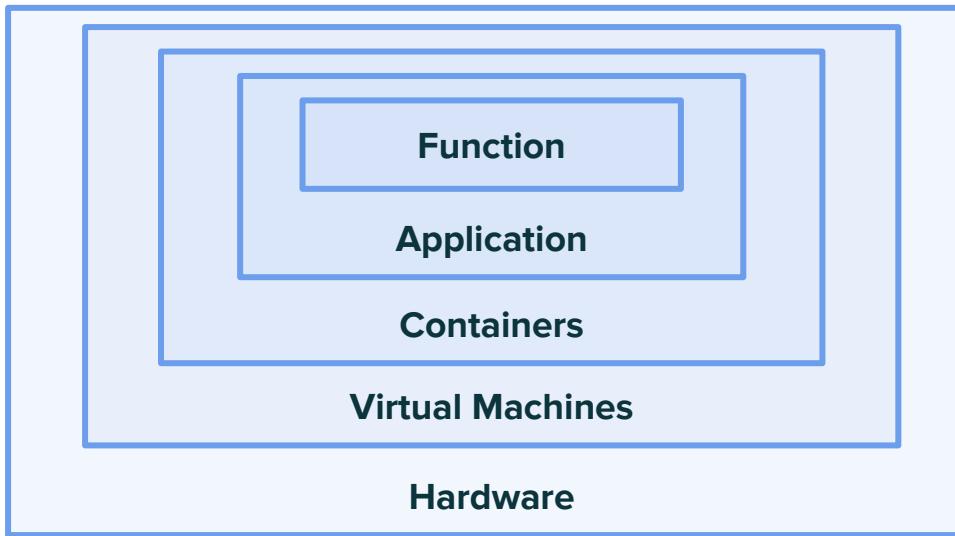
More Standards



Higher Efficiency



More Flexibility



Lower Efficiency

Coarse Grained Requirements

- Package to deployable
- Deploy and Run
- Map traffic and events to running handlers
- Conserve Resources

Finer Details

- Container creation and deployment automation
- Scale to zero and back
- Support revisions with traffic splitting and rollback
- Allows configurable event flows
- Gather metrics and telemetry from running deployments



Knative

Knative Codified Patterns

- Deploying a container
- Orchestrating source-to-URL workflows on Kubernetes
- Routing and managing traffic with blue/green deployment
- Automatically sizing workloads based on demand
- Binding running services to eventing ecosystems

Not a fully featured Function as a Service platform

Intended to be used by to build FaaS or PaaS

A screenshot of a Twitter post from Kelsey Hightower (@kelseyhightower). The post is a reply to @cpuguy83. The text reads: "Knative is not that top-level interface. That will come from the ecosystem. The containers and specs that describes how workloads run and flow between those systems is the value of Knative." The post was made at 7:29 AM - 8 Mar 2019.

Kelsey Hightower 
@kelseyhightower

Following

Replying to [@cpuguy83](#)

Knative is not that top-level interface. That will come from the ecosystem. The containers and specs that describes how workloads run and flow between those systems is the value of Knative.

7:29 AM - 8 Mar 2019

Knative is...

- An abstraction on top of Kubernetes
 - Deploys and runs containers
- Set of building blocks to construct your own FaaS/AaaS/CaaS
 - abstracts common tasks through custom Kubernetes API objects
- Middleware components essential to build modern, source-centric, and container-based applications that can run anywhere.

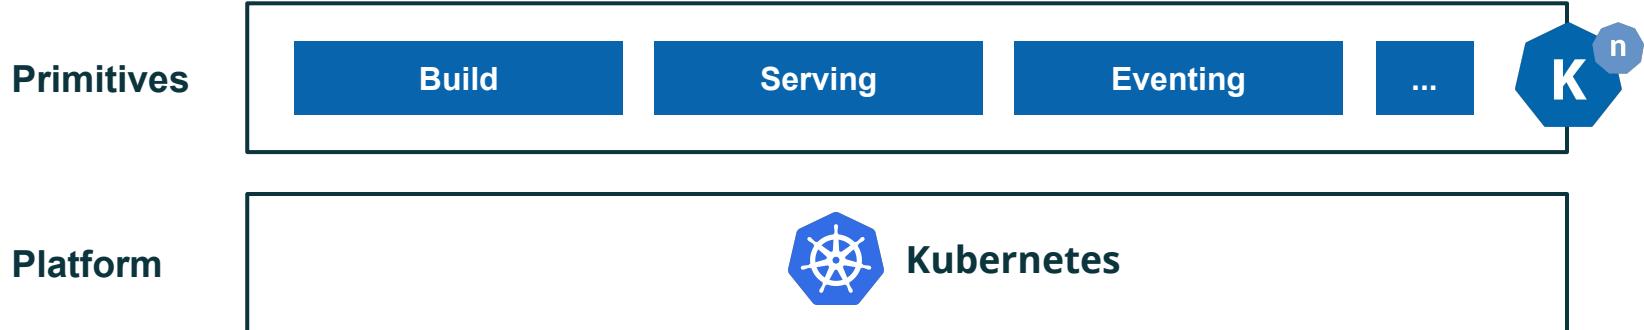
Knative architecture

Platform

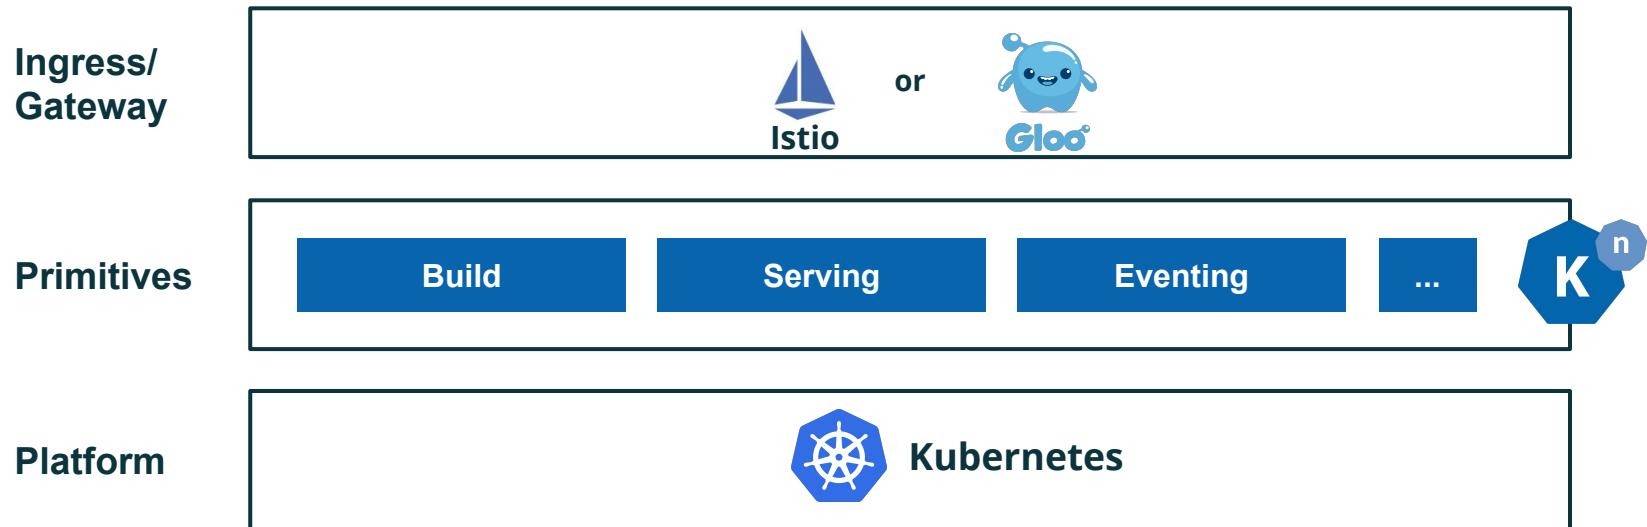


Kubernetes

Knative architecture



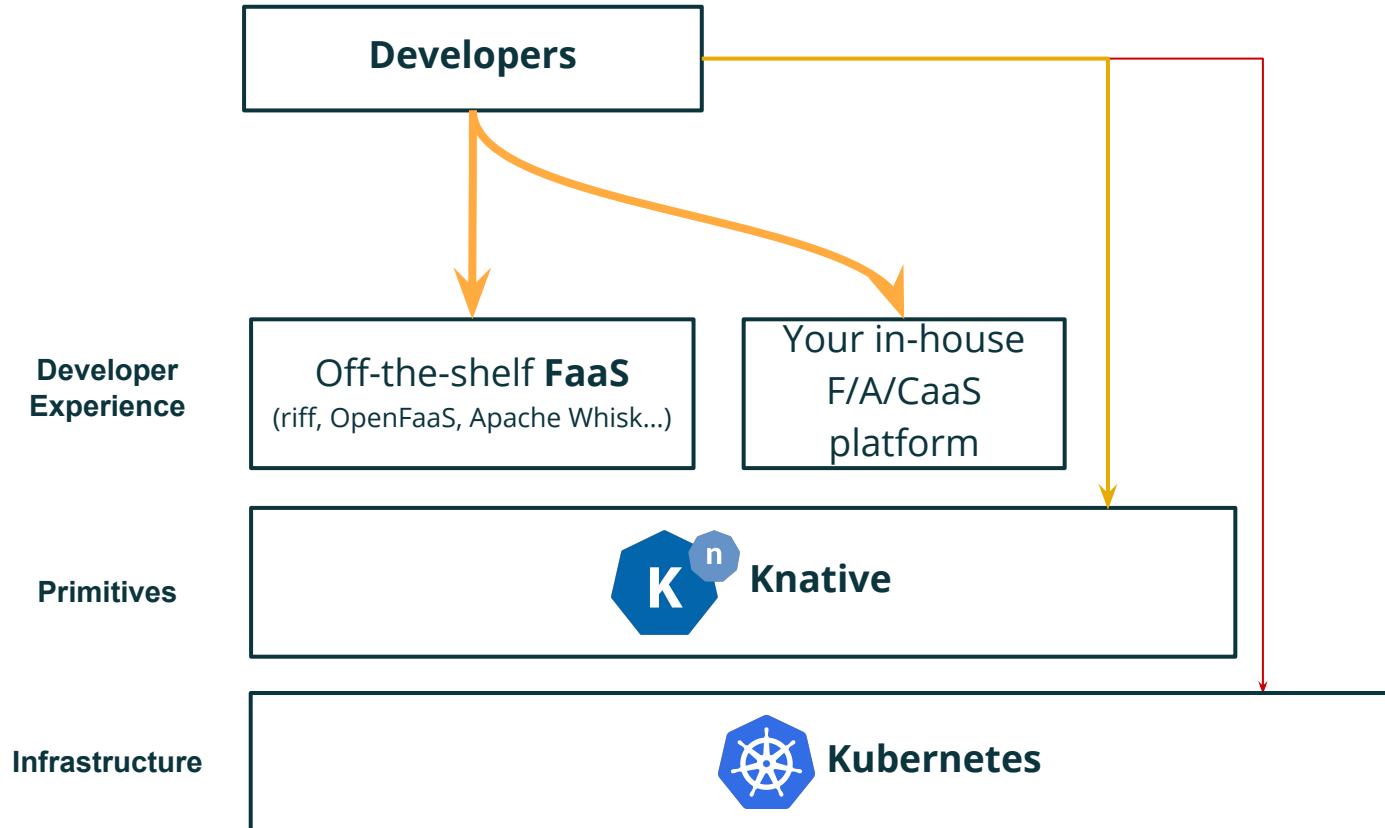
Knative architecture



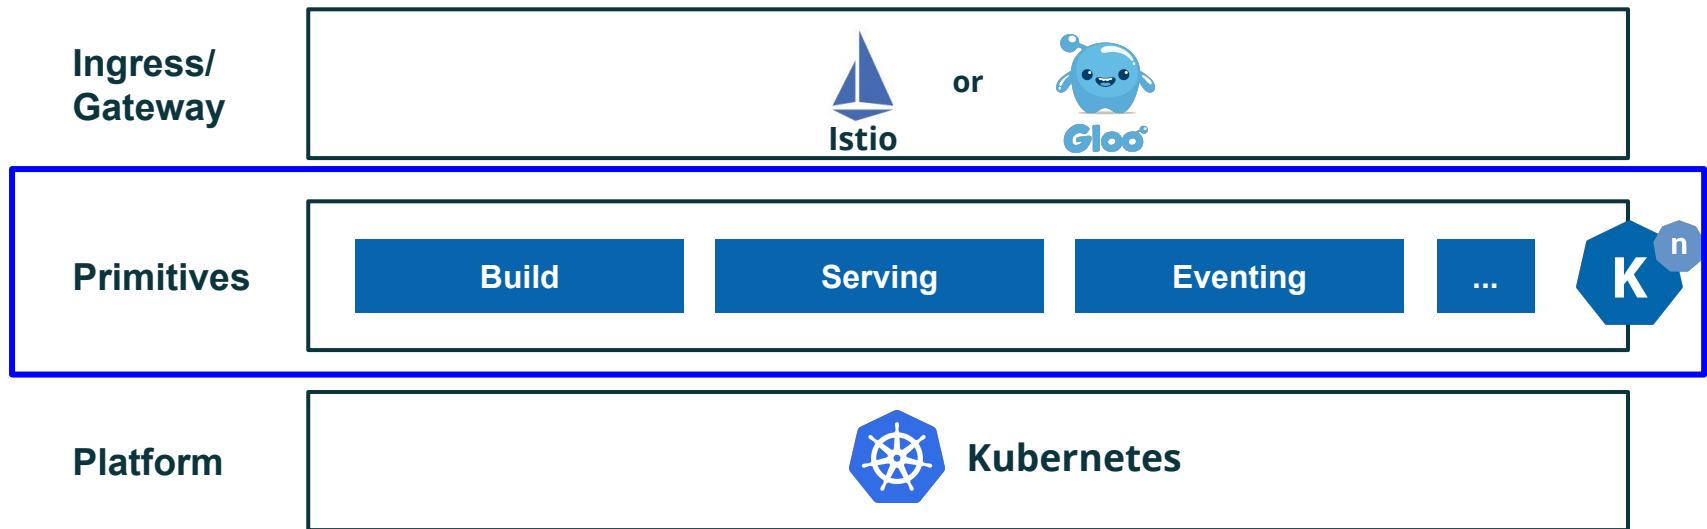
Audiences



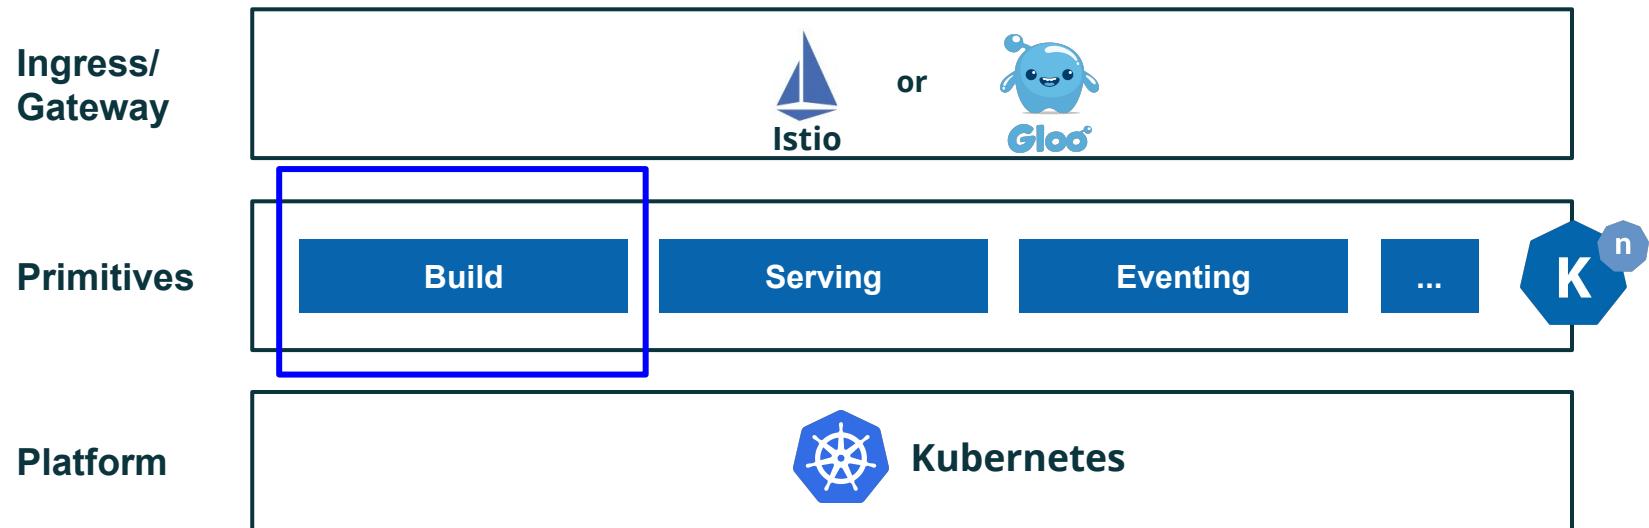
Developer Abstractions

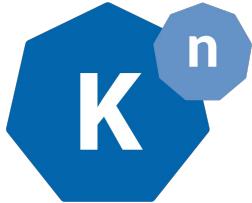


Knative Components



Knative Components





Build

Source-to-container build orchestration

Build

- **Defines process that runs to completion and provide status**
- **Create container image from supplied source for deployment to Knative Serving**
- **Connects to secured resources via Kubernetes Service Account**

Build Components

- Build is made up of Steps containing builders
- Builder - container image that accomplishes a task (discrete or composite)
 - One of the steps typically push to container registry

Build Components (Cont.)

- **Source - data necessary for build in mounted volume**
 - **GCS**
 - **Git**
 - **Custom (container image)**

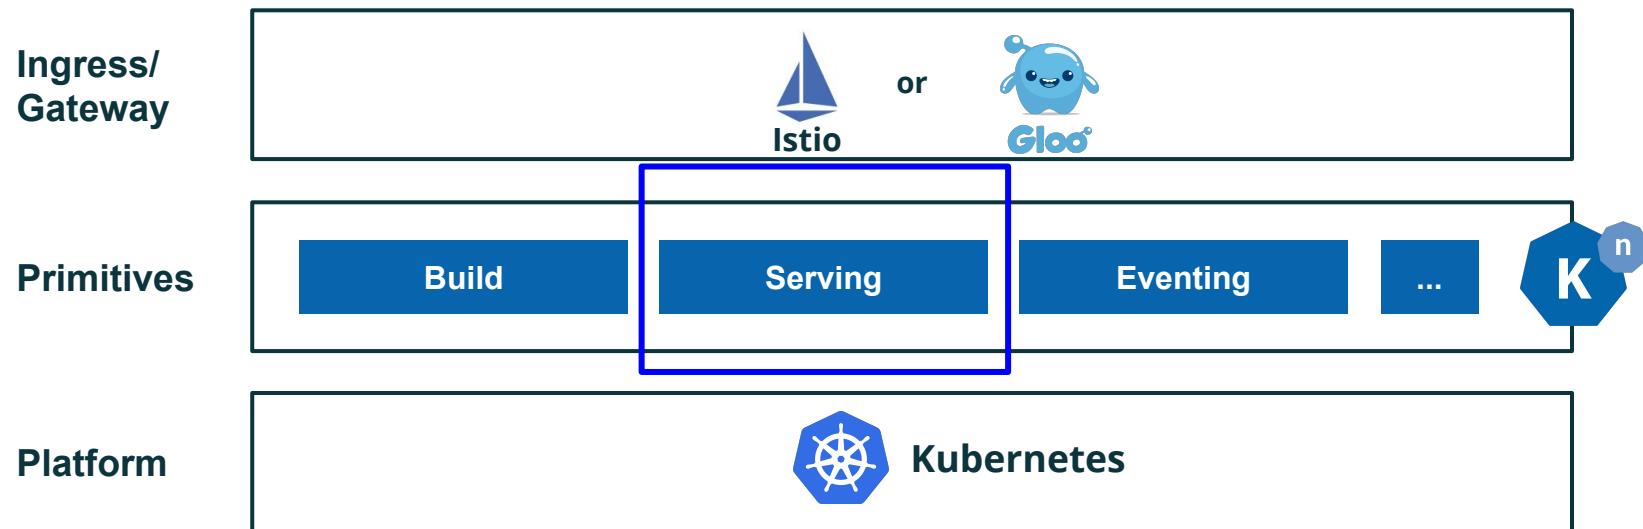
Build Components (Cont.)

- **Service Account** - privileges to run build
- **Volumes** - non default volumes to use during steps
- **Timeout** - max time for resource allocation and steps

Build Templates

- Encapsulate build process for reuse and consistency
 - limited parameterization
- Similar to “compile” phase of Heroku and Cloud Foundry buildpacks
- Limited to namespace but can be shared across cluster as Cluster Build Template

Knative Components





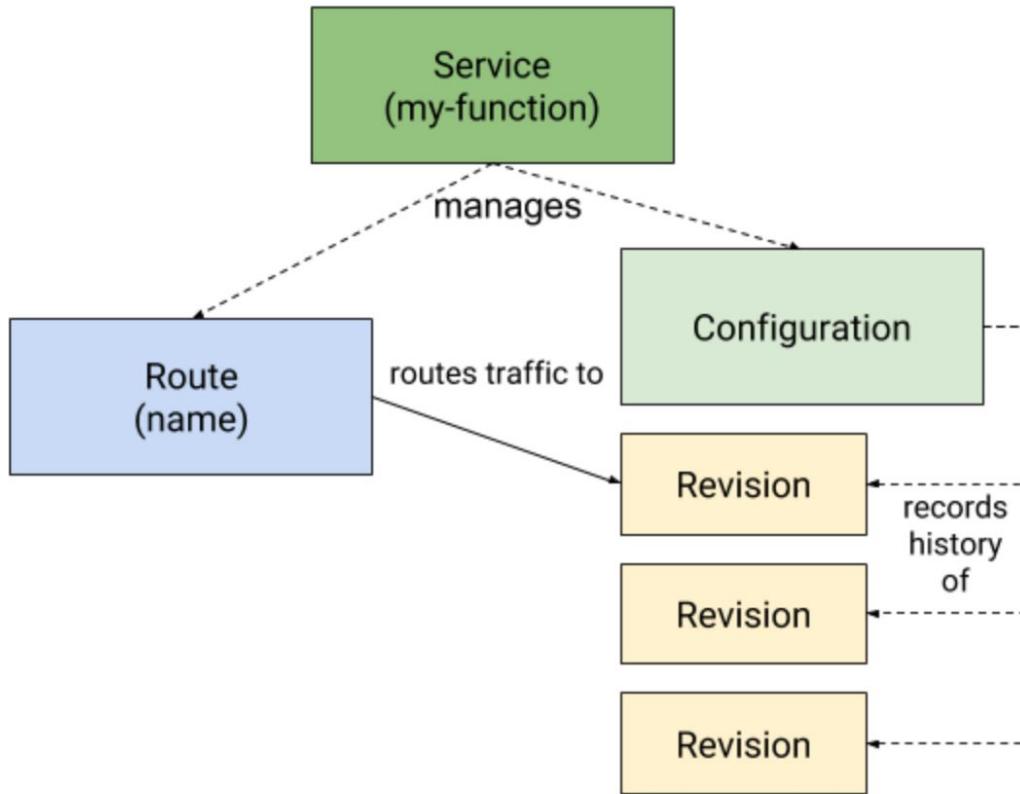
Serving

Request-driven serverless compute

Serving

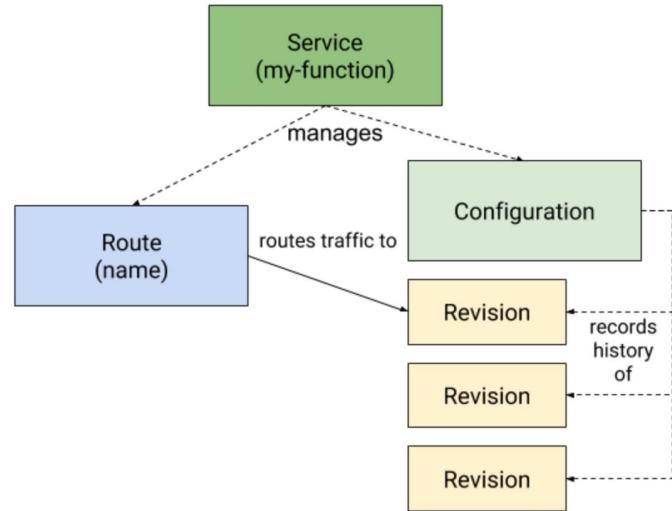
- **Rapid deployment of serverless containers**
- **Programming for ingress gateway for network routing (Istio or Gloo)**
- **Point-in-time snapshots of deployed code and configurations (enabling rollback)**
- **Automatic scaling up and down to zero**
- **Observability through logs and metrics**

Serving Components



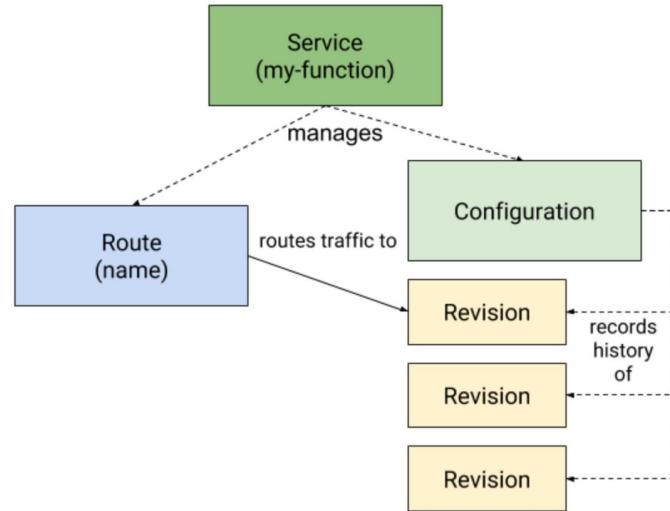
Route

- Maps HTTP routable endpoint to revision(s)
- Supports percentage based splitting of traffic



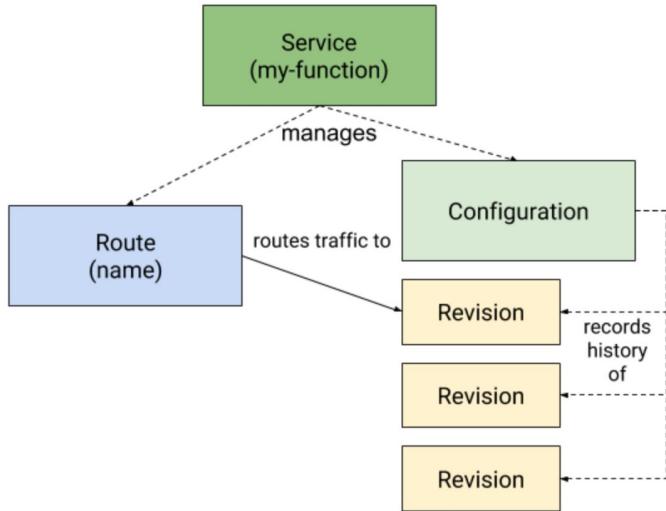
Revision

- **Immutable code (as container or build) and config**
- **Generated by configuration changes**
- **Retired no route mapped (resources deleted)**



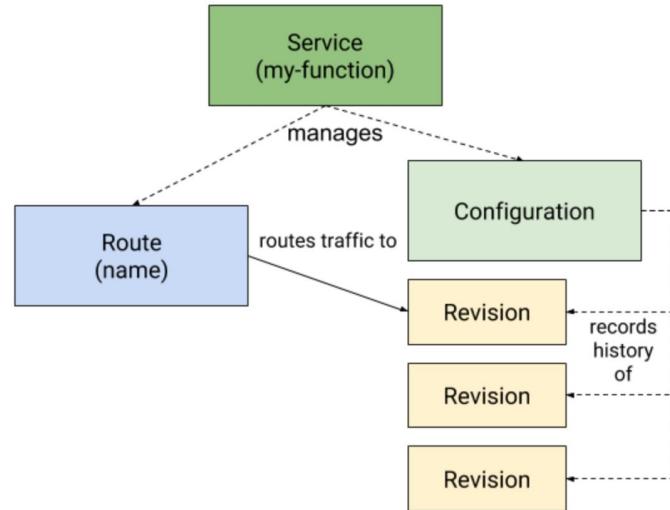
Configuration

- Latest revision of code and config



Service

- Set of routes and configurations/revisions
- Single abstraction of software component



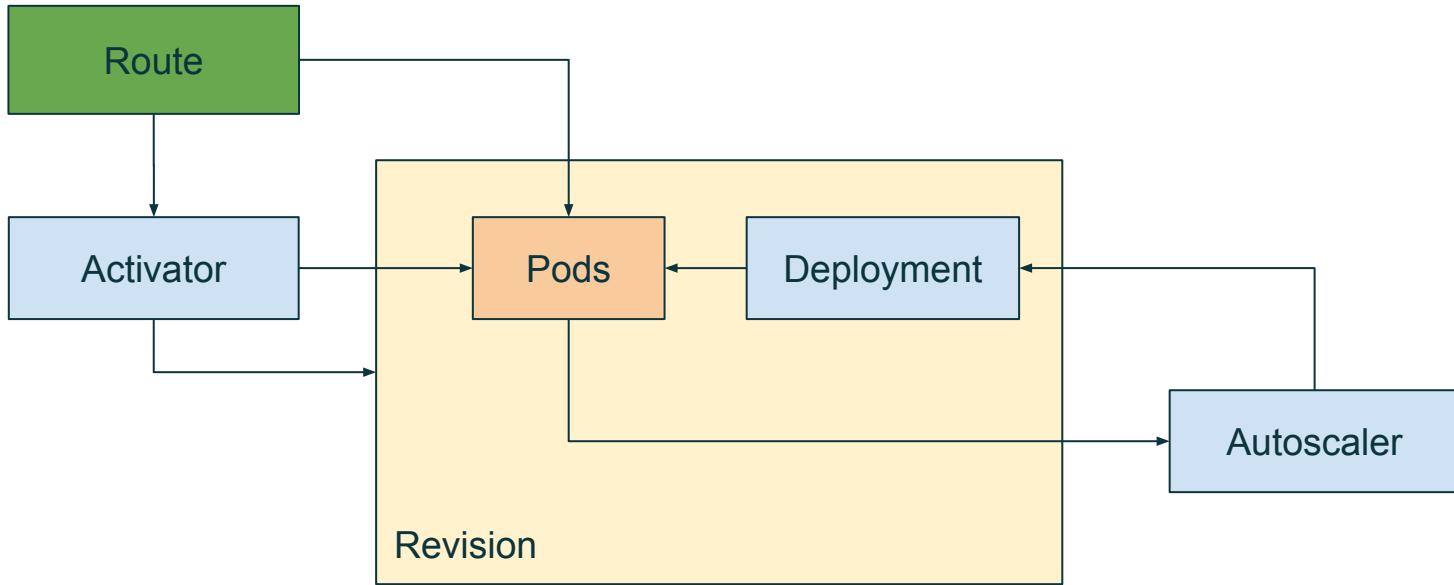
Autoscaling Goals

- **Fast - scale from 0 to 1000 concurrent requests in < 30 seconds**
- **Light - System makes decision without operator intervention**
- **Make Everything Better - Custom components are short term until K8s scaling components improved**

Pod States

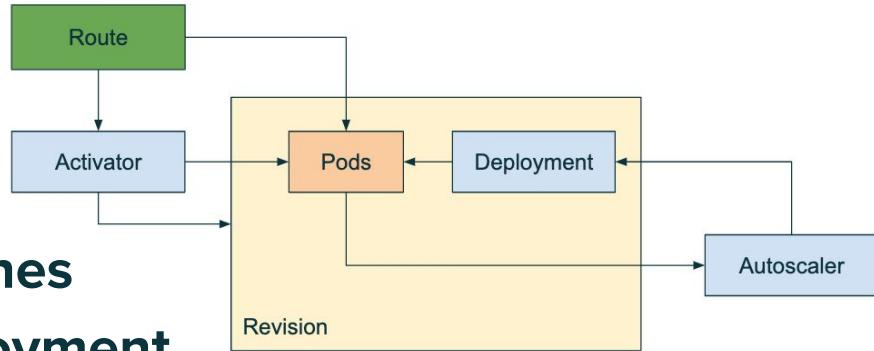
- **Active - Handling current traffic**
- **Reserved - Scaled to 0**
- **Retired - No route mapped**

Autoscaling Components



Autoscaling Components

- **Knative Deployment** - manages running revision
- **Knative Service Autoscaler** - watches revision metrics and advises deployment to scale increase/decrease pods
- **Knative Activator** - intercepts requests for revisions in reserved state, activates pods, forwards requests to activated pod



Autoscaling Modes

- **Stable - 60 second window averages**
 - calculated every 2 seconds
- **Panic - 6 second window averages**
 - Triggered when average is 2x desired
 - Increases only
 - Revert to stable after 60 seconds of stability
- **Deactivation - average of 0 for configured threshold**
 - Revision to reserved state
 - Default threshold 5 minutes
 - Minimum threshold 30 seconds

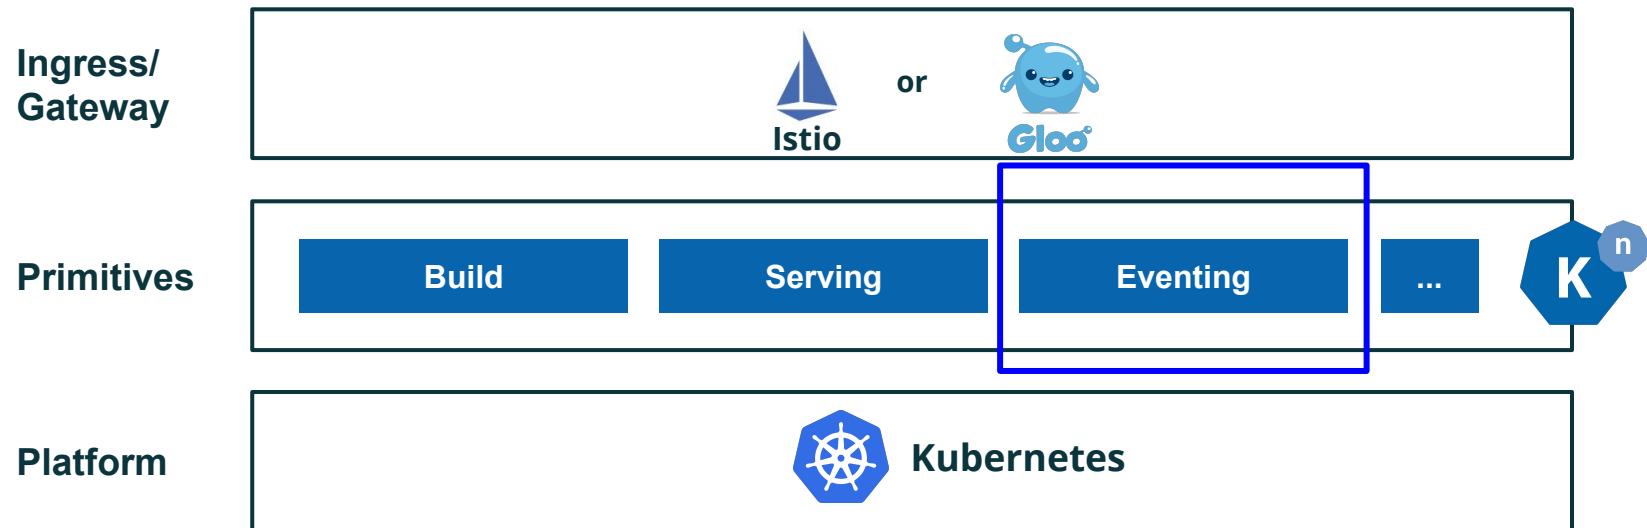
Autoscaling Classes

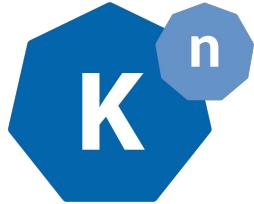
- **Knative Autoscaler - concurrent requests per pod**
 - default 100
- **Kubernetes Horizontal Pod Autoscaler - CPU Percent**
 - Default 80%

Logging And Metrics

- Knative collects logs using Fluentd
- Plugable logging/observability implementations
 - Elasticsearch and Kibana
 - Stackdriver
 - Custom logging plugin
- Provides Grafana metrics observability

Knative Components





Eventing

**Declarative management and delivery of
events**

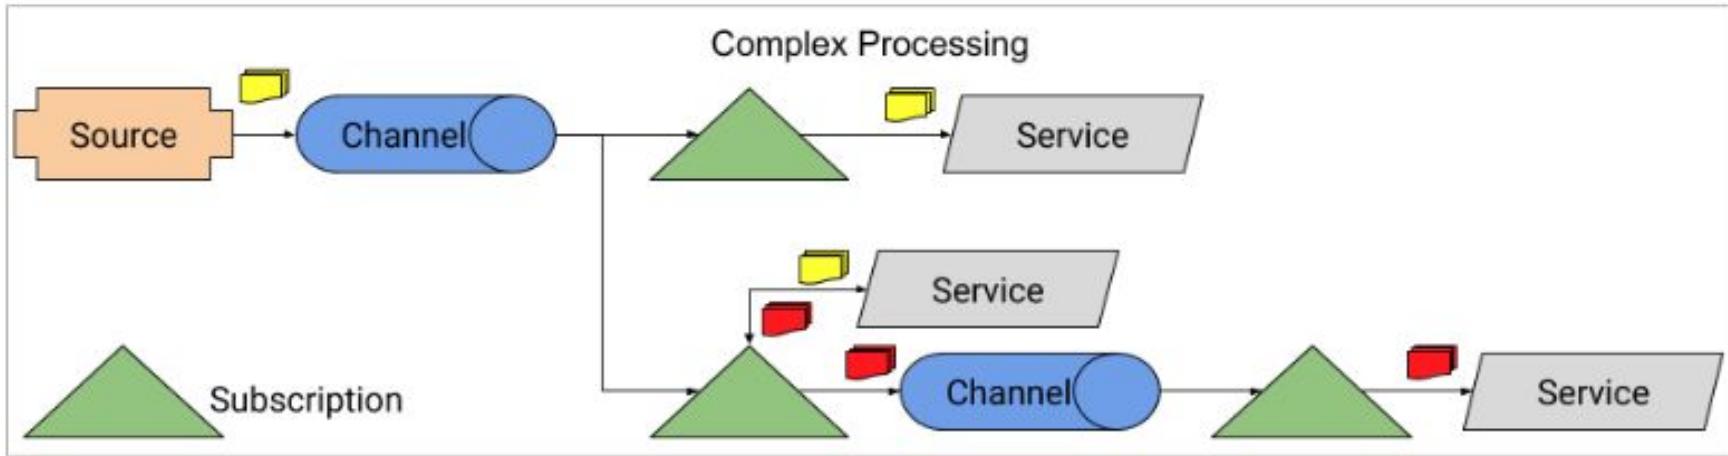
Eventing

- **Composable primitives**
- **Late-binding loosely coupled event producers and consumers**
- **Consistent with CNCF Serverless WG CloudEvents specification**

Services as Event Consumers

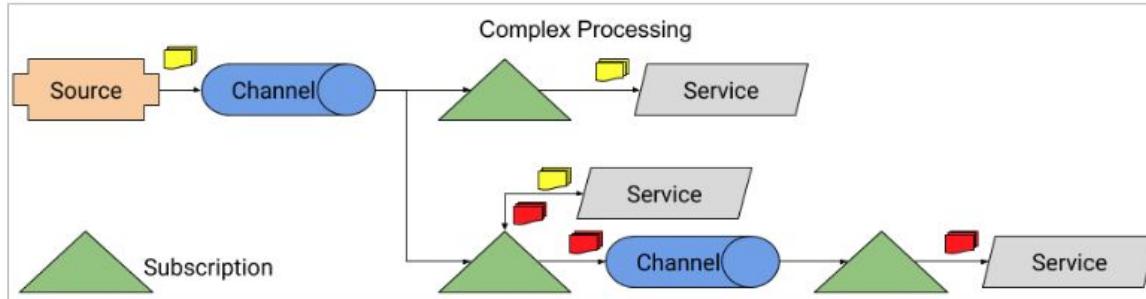
- Select and target specific subsets of the events from their producers
- Create new applications without modifying producer or consumer

Eventing Components



Source

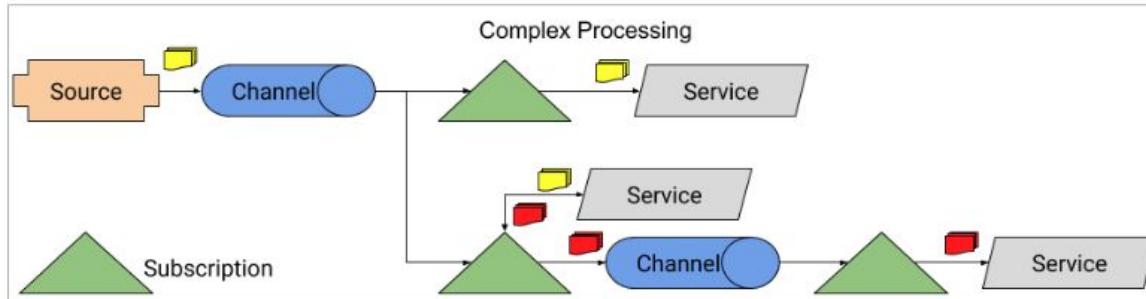
Register interest in events from other systems and publish to a sink (K8s events, Kafka, GitHub, GitLab, BitBucket, custom sources)



Service

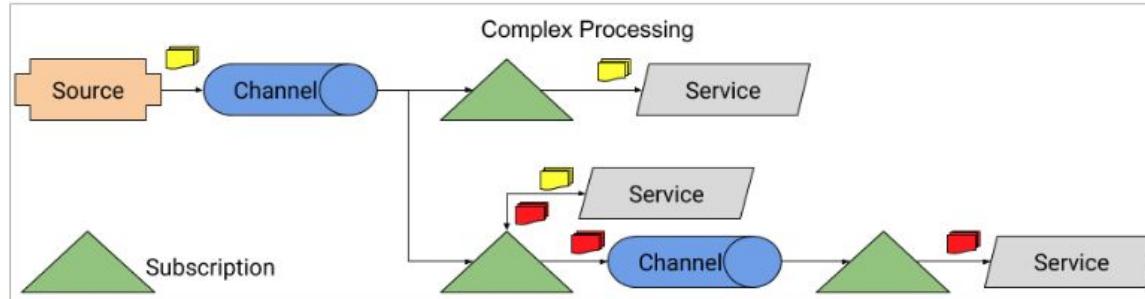
Consumer of events

- Addressable acknowledge delivery
- Callable - return 0 or 1 events



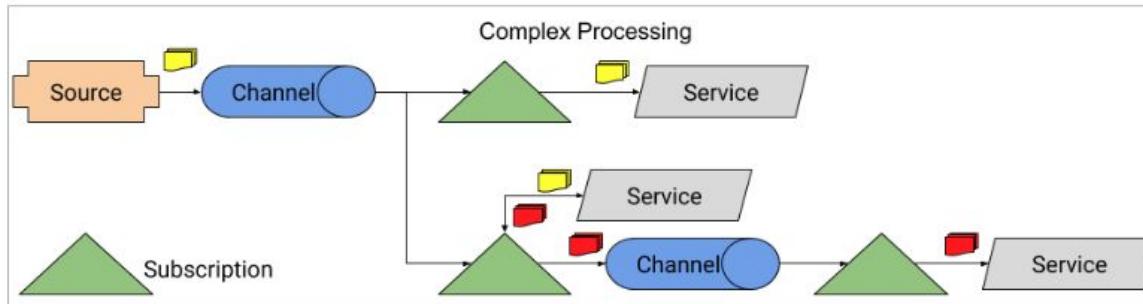
Channel

- Event forwarding persistence
- Plugable ClusterChannelProvisioner implementation for message providers

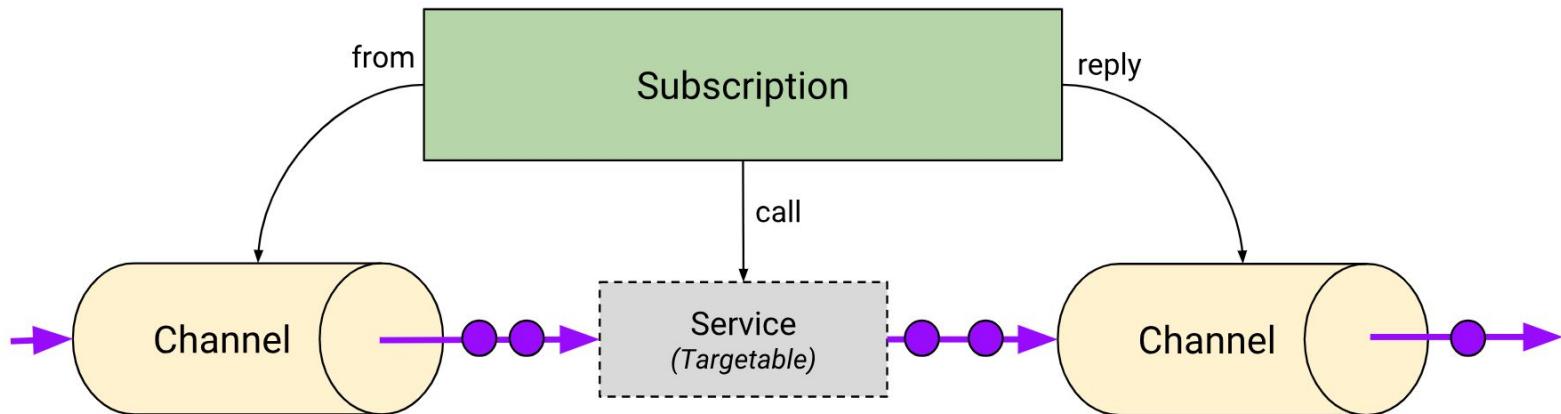


Subscription

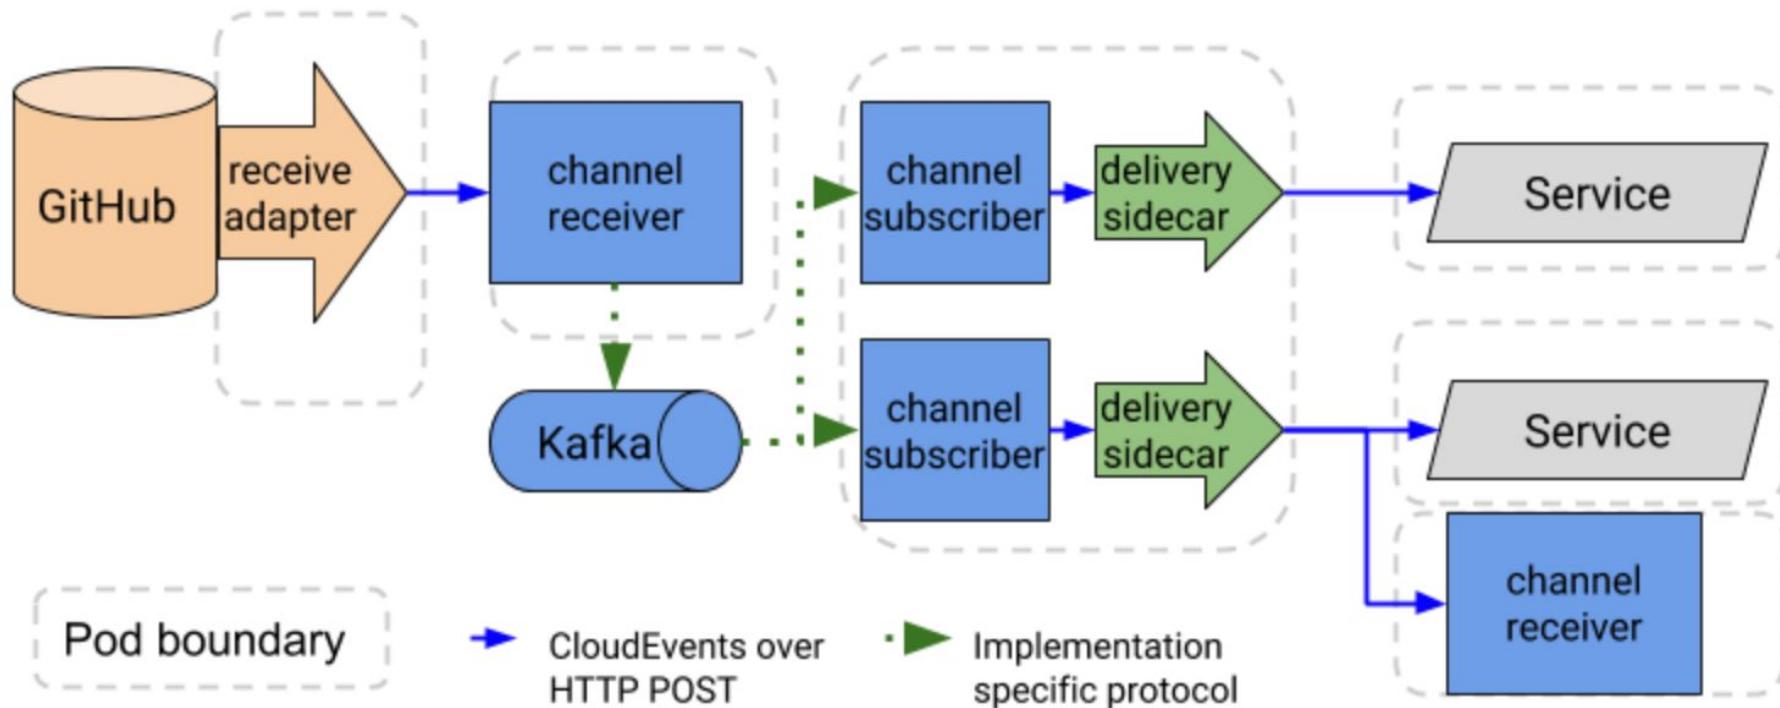
Means to declare interest in events from a channel



Channels & Subscriptions

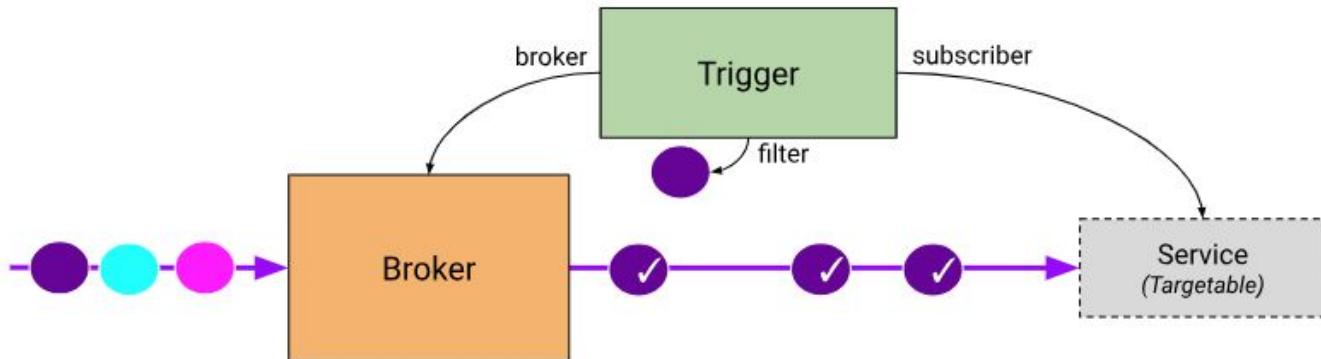


Data Plane



Eventing Components (New as of Knative 0.5)

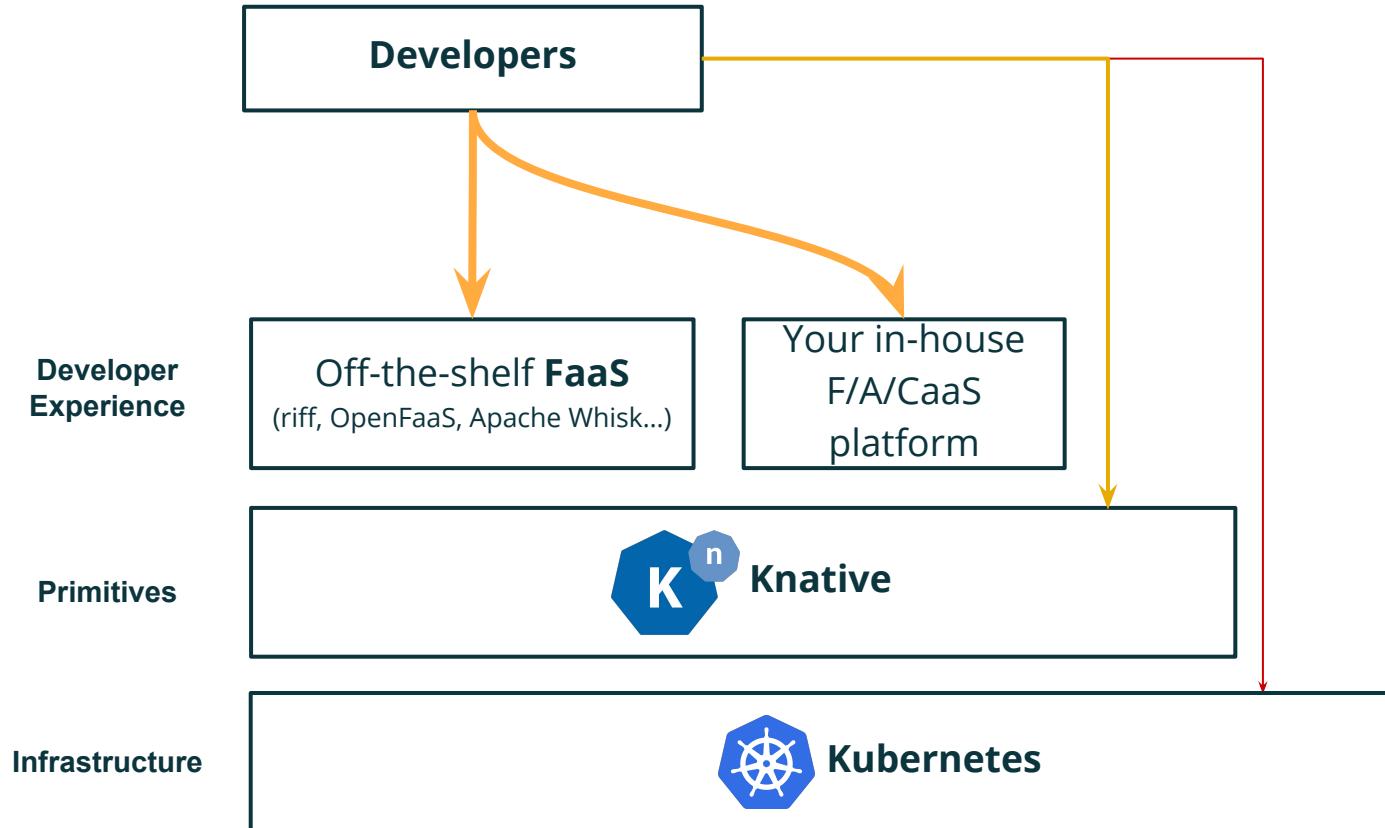
- **Broker** - mesh of events sent to interested subscribers
 - Use a channel template to create internal channels
- **Trigger** - declares subscription to events from a broker
 - Can be used to filter events based upon type



A large stack of red and tan clay bricks, some stacked vertically and others horizontally, creating a textured, layered appearance.

**Remember these are
primitives to build PaaS/FaaS**

Developer Abstractions





Questions???

Where to go to learn more

- <http://knative.dev>
- <https://github.com/knative/>
- <https://starkandwayne.com/blog/introduction-to-knative/>
- <https://medium.com/@pczarkowski/introduction-to-knative-b93a0b9aeeef>
- <https://www.ibm.com/cloud/learn/knative>
- <https://cloud.google.com/knative/>

Config Sample Links

- Build
 - Build - <https://github.com/GoogleCloudPlatform/knative-build-tutorials/blob/master/docker-build/README.md>
 - Build Template - <https://github.com/knative/build-templates/blob/master/kaniko/kaniko.yaml>
 - Build with template - <https://github.com/knative/build-templates/tree/master/kaniko>
- Serving
 - Deploy and routing examples -
<https://github.com/knative/docs/blob/master/docs/serving/samples/blue-green-deployment.md>
- Eventing
 - Channel - <https://github.com/gswk/knative-eventing-demo/blob/master/channel.yaml>
 - Subscription - <https://github.com/gswk/knative-eventing-demo/blob/master/subscription.yaml>
 - Broker - <https://github.com/knative/eventing/tree/master/docs/broker>
 - Trigger - <https://github.com/knative/docs/blob/master/docs/eventing/samples/gcp-pubsub-source/trigger.yaml>